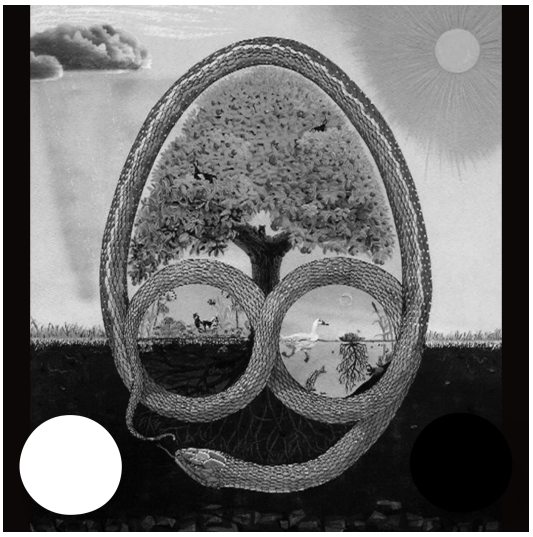

Unturf Automated General Intelligence: Merkle Providence Reverse RAG

Provenance-Preserving Cache Chains for Small Language Models

How Merkle trees turn Reverse RAG into a verifiable, shared knowledge layer.

How Hermes 3 Llama 3.1 8B answers with the confidence of a verified record, not a guess.

Russell Ballestrini <russell@unturf.com>, David Wong <david.s.wong@pm.me>, Riley Morgan <guybriley02@gmail.com>

unfirehose.com · uncloseai.com · permacomputer.com

April 2026

Abstract

License: AGPL-3.0-only · This algorithm, its implementation, & all associated code carry the GNU Affero General Public License v3.0 (only). You may use, modify, & distribute under those terms. No proprietary relicensing exists.

Reverse Retrieval Augmented Generation (Reverse RAG) demonstrated that client-side context injection lets small 8B-parameter language models outperform much larger models on page-specific questions. Our reference implementation runs **Hermes 3**, a NousResearch instruction fine-tune of Meta's Llama 3.1 8B, served freely at uncloseai.com. Our prior work showed that the client already holds the document. No vector database required. No embedding pipeline. No retrieval failures.

Merkle Providence Reverse RAG extends this with a second insight: the client not only holds the document, it can *remember what it already computed about that document*, prove that memory came from an unmodified source, & share that verified knowledge with others on public or private chains.

A Merkle tree over document content chunks produces a root hash: a 32-byte fingerprint that changes when any part of the document changes. Pairing this root hash with a question hash yields a cache key. A cache hit means a previously computed answer returns instantly, with a Merkle proof of origin. A cache miss triggers fresh inference, stores the result, & extends our chain.

Unfirehose serves as our chain substrate. Every inference session produces JSONL records that feed into SQLite at `~/.unfirehose/unfirehose.db`. Public chains broadcast verified answer records to peers. Private chains stay local. Either way, a small model gains access to a growing body of pre-verified answers, each carrying cryptographic proof of the document version that produced it.

This combination removes two limits that have constrained small models:

1. **The repetition tax**: answering the same question about the same document on every visit, burning inference tokens & time, producing answers that may vary slightly each run.
2. **The trust deficit**: small models hallucinate. Cached answers with Merkle proofs do not hallucinate. They either match a verified record or they do not exist yet.

Once published, this technique forces a reckoning: every RAG system that cannot prove provenance of its retrieved chunks operates on unverifiable context. Merkle Providence makes provenance a first-class primitive.

We call this stack **Web 2.5**: the existing web's documents & URIs, augmented with machine learning context injection & Merkle-verified answer chains, without requiring a blockchain, a new protocol, or replacement of any existing infrastructure. Web 2.5 runs on top of what already exists. Any page. Any browser. Any small model. Public or private chain. No permission required.

Reference Model: Hermes 3 Llama 3.1 8B

Our reference implementation runs on **Hermes 3**, a fine-tune of Meta's Llama 3.1 8B developed by NousResearch. Hermes 3 builds on Llama 3.1's 128k context window & adds instruction alignment, function-calling discipline, & reasoning consistency that raw base models lack. NousResearch trained Hermes 3 to follow multi-turn conversation formats precisely, making it reliable for our structured providence cache interactions.

We chose Hermes 3 Llama 3.1 8B for three reasons:

Size matches our thesis. 8 billion parameters represent a model any person can run on consumer hardware. Our reference deployment hardware & quantization details stay current at uncloseai.com/inference.html. Our technique produces no value if it requires expensive hardware to demonstrate.

Fine-tuning demonstrates our point. Hermes 3 shows that fine-tuning a base model for instruction following, not scaling parameters, produces the reliability gains that matter for user-facing applications. Our Merkle Providence layer extends this logic: fine-tuned behavior + verified cached context beats raw parameter count.

uncloseai.com serves it freely. Our public inference endpoint at uncloseai.com runs Hermes 3 Llama 3.1 8B at no cost to any user. Our Reverse RAG pipeline in `uncloseai.js` targets this endpoint by default. Our Merkle Providence cache layer builds directly on this existing infrastructure. Any person anywhere, on any device with a browser, accesses our same Hermes 3 instance our paper describes.

Hermes 3 Llama 3.1 8B outperforms many larger models on page-specific questions when Reverse RAG supplies full context. With Merkle Providence caching, it delivers consistent, verifiable answers across sessions. Our paper uses Hermes 3 throughout all examples & benchmarks.

1. The Problem Reverse RAG Left Open

Reverse RAG solved context retrieval. Client-side DOM extraction replaced vector databases. Full-page injection replaced chunk fragmentation. Small models with perfect context outperformed large models with partial context.

One problem remained: **repetition**.

A user visits a technical document. Hermes 3 Llama 3.1 8B reads 8,000 tokens of page content, classifies the page, & answers their question. One hour later, a different user visits our same document. Our same Hermes 3 instance reads our same 8,000 tokens, classifies our same page, & answers a similar question again. Tomorrow, a third user. Our same 8,000 tokens. Our same classification. Our same inference cycle.

Nothing remembered. Nothing reused. No proof that any two answers came from our same source.

This repetition carries three costs:

- **Compute cost**: every visit re-runs inference on content that has not changed.

- **Consistency cost**: stochastic inference on identical inputs produces slightly different answers. Two users asking near-identical questions may receive meaningfully different responses from our same page.
- **Trust cost**: no mechanism exists to verify that an answer actually derived from a specific document version. A model can claim anything.

Merkle Providence addresses all three.

2. Merkle Trees as Content-Addressed Cache Keys

A Merkle tree splits a document into chunks, hashes each chunk, then builds a binary tree of hashes upward until a single root hash remains. This root hash has two properties that make it ideal for caching:

Property 1: Determinism. Our same document always produces our same root hash. Two clients processing our same page independently compute identical roots. No coordination required.

Property 2: Tamper sensitivity. Any change to any byte of the document changes at least one leaf hash, which propagates upward to invalidate our root. A document that changed since our last visit produces a new root hash automatically. Our cache never serves a stale answer for a modified document.

Cache key construction:

```
document_root = merkle_root(chunks(page_content))
question_key  = sha256(normalize(user_question))
cache_key     = document_root + ":" + question_key
```

A cache hit at this key means: *an earlier session computed an answer to this question, about this exact version of this document.* Our answer arrives in O(1) time. Our Merkle proof confirms our document version.

A cache miss means: run inference, store our result, extend our chain.

Git & Mercurial are already Merkle trees. For source code repositories, we do not need to build our own Merkle tree. Every git commit hash IS a Merkle root: a SHA-1 (or SHA-256 in git 2.x) digest of our entire repository tree at that point in time. Every hg commit carries an identical guarantee. The version control system has already done our hashing work.

For codebases, our cache key simplifies to:

```
document_root = git_commit_hash   # or hg changeset hash
question_key  = sha256(normalize(user_question))
cache_key     = document_root + ":" + question_key
```

When a developer commits, our cache key changes automatically. Prior answers expire. Fresh inference runs against our new commit. No manual invalidation. No polling. Our version control system signals our cache boundary by design.

3. The Providence Layer: Proof of Origin

"Providence" carries two meanings that both apply here:

Provenance: where something came from. A Merkle proof lets any verifier confirm that an answer derived from a specific document version without receiving our full document. Our proof path, a sequence of sibling hashes from leaf to root, constitutes a compact cryptographic certificate of origin.

Foresight: the system appears to "already know" answers to questions asked before. From a user's perspective, a cache hit looks like instant comprehension. Our model responded before seemingly finishing to read. This resembles foresight because our prior sessions did our work ahead of time.

A providence record for a cached answer contains the full inference context:

```
{
  "cache_key":        "abc123...:def456...",
  "document_root":    "sha256:abc123...",
  "document_uri":     "https://example.com/docs/api",
  "question_hash":    "def456...",
  "question_text":    "What does the rate_limit field mean?",
  "answer_text":      "The rate_limit field specifies...",
  "merkle_proof":     ["sha256:111...", "sha256:222...", ...],
  "model_id":         "adamo1139/Hermes-3-Llama-3.1-8B-FP8-Dynamic",
  "base_uri":         "https://hermes.ai.unturf.com/v1",
  "temperature":      0.1,
  "source_type":      "web",
  "git_commit":       null,
  "conversation_hash": "9a3f...",
  "turn_number":      1,
  "created_at":       1744000000,
  "hit_count":        0
}
```

Our providence record collects maximum metadata, but only a subset enters our cache key. This distinction matters.

Tier 1 — Cache Key Inputs. A difference in any of these fields produces a distinct cache entry:

- `document_root` — content fingerprint (Merkle hash or git/hg commit)
- `question_hash` — SHA-256 of normalized question text. Four-step canonicalization, in order: NFC + whitespace collapse + edge-strip, then case-fold to lowercase, then trailing-punctuation strip, then drop standalone English articles (`the` / `a` / `an`). Trailing chars in the strip set today: `.?!,;:` (ASCII), ■■■■ (CJK full-width), … (ellipsis). Pairs (`" ' ) ] }`) are NOT stripped — naive one-sided stripping breaks balance, and apostrophes carry meaning (`X's` ≠ `X`). Articles are matched as exact lowercase standalone tokens, so substrings like `thesis` (contains `the`) stay untouched. Equivalence class: `who is X`, `who is X?`, `Who Is X.`, `who is the X`, `who is a X`, `who is an X■` all map to the same `question_hash`, so a providence_cache hit returns the same answer for any of those forms.
- `model_id` — canonical model identifier
- `model_revision` — exact weights revision (HuggingFace commit hash or equivalent)
- `quantization` — fp16, bf16, fp8, int8, int4, q4_k_m, etc. Same weights at different quantizations produce measurably different answers.
- `conversation_hash` — SHA-256 of the full normalized industry-standard chat messages array: every turn in order, system prompt through the final user message. A single-turn Q&A and a 6-turn conversation that arrives at the same final question produce different answers & different cache keys. Callers hash client-side & send only the hash. Message content never enters our logs — Operation Voyeur.
- `seed` — RNG seed when set. A seeded run produces a deterministic answer; a different seed produces a different one. Null seed entries cache on the other five fields alone.

Tier 2 — Metadata Only. Stored for research, routing, & audit. Not part of the key:

- `base_uri` — inference endpoint URI
- `temperature`, `top_p`, `top_k` — sampling parameters
- `repetition_penalty`, `frequency_penalty`, `presence_penalty`
- `max_tokens`, `context_window`
- `backend` — vllm, llama.cpp, ollama, transformers, TGI
- `node_id` — mesh node that served the request
- `inference_ms` — wall-clock inference time

Tier 2 metadata does not affect cache correctness. Two answers produced at different temperatures from the same model, revision, & quantization, answering the same question about the same document, represent the same logical answer — stochastically equivalent. We store the temperature of the first run that populated our cache. Future hits return that answer regardless of the caller's temperature setting.

Any recipient of this record can verify our Merkle proof against our document root. If our document root matches what they compute from our same URI, our answer carries full provenance. If their computation produces a different root, our document changed, our answer no longer applies to their version.

4. Unfirehose as Chain Substrate

Unfirehose collects JSONL session data from machine learning harnesses across a mesh of compute nodes. Every tool call, every model response, every session event flows into `~/.unfirehose/unfirehose.db` via our background ingestion worker.

Our Merkle providence layer extends this existing infrastructure with a providence cache table:

```
CREATE TABLE providence_cache (
  cache_key        TEXT PRIMARY KEY,
  document_root    TEXT NOT NULL,
  document_uri     TEXT NOT NULL,
  question_hash    TEXT NOT NULL,
  question_text    TEXT NOT NULL,
  answer_text      TEXT NOT NULL,
  merkle_proof     TEXT NOT NULL,  -- JSON array of sibling hashes
  model            TEXT NOT NULL,
  created_at       INTEGER NOT NULL,
  chain            TEXT NOT NULL DEFAULT 'private',
  peer_count       INTEGER NOT NULL DEFAULT 0
);

CREATE INDEX idx_providence_root ON providence_cache(document_root);
CREATE INDEX idx_providence_uri  ON providence_cache(document_uri);
```

Two chain modes:

Private chain: records stay local at `~/.unfirehose/`. Our personal session history. Our model learns from our own prior work on documents we visit repeatedly. Our chain belongs to us alone.

Public chain: records broadcast to peers via our unfirehose router daemon. Any node on our mesh that has processed our same document root can contribute answers. Any node that queries our same document root can receive pre-verified answers from others. Our chain grows with every participant.

Public chains do not share document content. They share proofs. Our Merkle proof proves an answer came from a specific document version without transmitting our document. Document content stays on our originating client, exactly as Reverse RAG intended.

5. How Small Models Punch Outside Their Weight Class

Parameter count governs a model's ability to reason from first principles. Context quality governs a model's ability to answer questions about specific documents. Verified memory governs a model's ability to give consistent, trustworthy answers over time.

Large models win on reasoning from first principles. Small models, given the right tools, win on everything else.

Merkle Providence Reverse RAG gives small models three tools:

Tool 1: Full context (from Reverse RAG). Our same page injection described in our prior paper. Our model receives every word, every heading, every entity, every code block of our document. Nothing retrieved by cosine similarity. Nothing fragmented. Our full document, live from our DOM.

Tool 2: Verified recall (from our providence cache). Before our model reads our document, we check our cache. A cache hit returns our answer in milliseconds with a Merkle proof. Our model delivers a verified

answer faster than any large model can load our full context. Our small model "already knew" because our prior session did our work.

Tool 3: Peer-verified answers (from public chains). On a public chain, our small model benefits from every session any participant ran against our same document version. A document processed by 100 sessions across our mesh produces a rich cache. Our Hermes 3 Llama 3.1 8B instance answers from our collective record, not just its own prior work.

These three tools operate at different timescales:

- **O(1)**: cache hit, local or peer. Instant. No inference required.
- **O(context)**: Reverse RAG with full page injection. Fast. One inference pass over complete context.
- **O(reasoning)**: fresh inference on novel questions. Standard. One inference pass, uncached.

A large model operates exclusively at O(reasoning). Our small model with Merkle Providence operates at O(1) for previously answered questions, O(context) for new questions on known documents, & O(reasoning) only for truly novel questions on new documents.

6. Architecture

Stage 1: Document Fingerprinting

Our same DOM extraction from Reverse RAG runs first. Once our extraction completes, we split our content into 512-token chunks & build our Merkle tree:

```
function merkleRoot(chunks) {
    let layer = chunks.map(c => sha256(c));
    while (layer.length > 1) {
        const next = [];
        for (let i = 0; i < layer.length; i += 2) {
            const left  = layer[i];
            const right = layer[i + 1] ?? left;  // odd node doubles itself
            next.push(sha256(left + right));
        }
        layer = next;
    }
    return layer[0];
}
```

Our root hash serves as our document fingerprint for our entire session. One hash computation. No server calls.

Stage 2: Cache Lookup

Before any LLM call, we check our providence cache:

```
async function lookupProvidence(documentRoot, questionText) {
    const questionHash = sha256(normalize(questionText));
    const cacheKey     = `${documentRoot}:${questionHash}`;

    // Check local chain first
    const local = await db.get(
        'SELECT * FROM providence_cache WHERE cache_key = ?',
        [cacheKey]
    );
    if (local) return { answer: local.answer_text, proof: local.merkle_proof, source: 'local' };

    // Check public chain peers
    if (publicChainEnabled) {
        const peer = await queryPeers(cacheKey);
        if (peer && verifyMerkleProof(peer.merkle_proof, documentRoot)) {
            await storeLocal(peer);  // cache locally for future hits
```

```
            return { answer: peer.answer_text, proof: peer.merkle_proof, source: 'peer' };
        }
    }

    return null;  // cache miss, proceed to inference
}
```

Stage 3: Inference & Cache Write

On a cache miss, our standard Reverse RAG inference runs. Once our answer arrives, we store our providence record:

```
async function storeProvidence(documentRoot, documentUri, question, answer, model) {
    const questionHash = sha256(normalize(question));
    const cacheKey     = `${documentRoot}:${questionHash}`;
    const merkleProof  = computeProof(documentRoot, question);  // path to root

    await db.run(`
        INSERT OR REPLACE INTO providence_cache
        (cache_key, document_root, document_uri, question_hash, question_text,
         answer_text, merkle_proof, model, created_at, chain)
        VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?, ?)
    `, [cacheKey, documentRoot, documentUri, questionHash, question,
        answer, JSON.stringify(merkleProof), model, Date.now(), chainMode]);

    if (chainMode === 'public') broadcastToMesh(cacheKey);
}
```

Stage 4: Proof Verification

Any party that receives a providence record can verify it:

```
function verifyMerkleProof(proof, claimedRoot, leafContent) {
    let current = sha256(leafContent);
    for (const sibling of proof) {
        current = sha256(current < sibling
            ? current + sibling
            : sibling + current);
    }
    return current === claimedRoot;
}
```

If our document root computed from our verifier's live DOM matches `claimedRoot`, our answer carries full provenance. Our answer derives from our exact document version our verifier currently views. If our roots diverge, our document changed. Our cache entry no longer applies.

7. Public vs. Private Chains

Property	Private Chain	Public Chain
Storage	`~/.unfirehose/` local only	Broadcast via unfirehose router mesh
Participants	One user, one machine	Any node on our mesh
Document content shared	Never	Never (proofs only)
Answer content shared	Never	Yes, with Merkle proof
Trust model	Self-trust: our own prior work	Peer-trust: verified by Merkle proof
Best for	Personal document libraries, sensitive internal docs	Open documentation, public research, shared reference material

Cache growth rate	Linear in our own sessions	Linear in mesh participant sessions

Private chains suit corporate knowledge bases, personal research notes, any document our user does not want to share. Our local model benefits from our own prior sessions. No data leaves our machine.

Public chains suit open documentation, public APIs, research papers, any document where sharing verified answers benefits our community. No document content leaves our originating client. Proofs travel instead.

In our proposed design, our unfirehose router daemon would manage chain membership & peer discovery. Our user opts in to public chains per-document-domain or per-session. Our current implementation stays private.

8. Forcing the Issue

Some technical formulations compel structural change not because they require adoption but because their existence makes our prior approach indefensible.

Our Merkle Providence layer forces one question onto every RAG system: *can you prove where your context came from?*

Traditional RAG retrieves chunks by cosine similarity. No chunk carries a Merkle proof. No retrieval carries a document fingerprint. Our retrieved context could come from a stale index, a corrupted chunk, or a document that changed since our last ingestion. Our RAG pipeline trusts our embedding similarity as a proxy for relevance. Nothing verifies our source.

Once Merkle-proved answers exist in our ecosystem, unverified RAG answers carry an implicit disclaimer: *I retrieved this from somewhere, I believe it matches your question, but I cannot prove our source was our document you currently view.*

Our providence cache answers carry a different posture: *I computed this answer from our exact bytes of our document whose root hash I can show you. Verify it yourself.*

This posture change does not require every system to adopt Merkle trees immediately. It requires every system to acknowledge that provenance matters, & that systems which cannot demonstrate provenance operate at a disadvantage in trust-sensitive contexts.

Medical documentation. Legal reference material. Security advisories. Financial disclosures. Any domain where "I retrieved something similar" falls short of "I can prove this came from our exact source you cited."

9. What the Providence Cache Curates & Aggregates

Every inference pass through our Merkle Providence layer produces a verified record. Over time these records accumulate into a structured, queryable knowledge base. Here we describe what our cache collects & what emerges from that collection.

Document fingerprints. Every document processed — web page, git-versioned file, or raw content — contributes a `document_root`: a deterministic hash of its content at a specific moment. Our cache accumulates a catalog of every document any session has processed, keyed by content not by name or URI. A document that moves keeps its fingerprint. A document that changes gets a new one.

Answer records. For every (document, conversation, model, revision, quantization) combination, our cache stores one verified answer. The answer record carries: the question text, the answer text, the full Merkle proof path, the conversation hash that produced the context, & the model identity & inference configuration that produced the output. These records aggregate into a corpus of verified knowledge about our document collection.

Conversation context fingerprints. Our `conversation_hash` field captures the full industry-standard chat messages array, system prompt, all prior turns, & the final user message, as a single SHA-256. Our cache accumulates not just answers but the exact conversational paths that produced them. Two agents that reached the same question via different conversation histories leave distinct fingerprints. Over time we observe which conversational approaches yield cache hits & which always miss — signal for prompt engineering.

Hit counts & reuse rates. Every cache lookup that finds an existing record increments `hit_count` & updates `last_hit_at`. Our cache aggregates these into a measure of knowledge reuse: which documents get queried repeatedly, which questions recur, which model + quantization combinations serve the most requests. High-hit records represent our most valuable verified knowledge — answers our mesh has found useful enough to retrieve many times.

Model & endpoint diversity. Each record carries `model_id`, `model_revision`, `quantization`, & `base_uri`. Our cache aggregates a record of which models answered which questions about which documents. Over time this produces a comparative dataset: the same question about the same document answered by Hermes 3 at fp8, at int4, by a different model entirely. Our cache does not collapse these — it preserves the full diversity for analysis.

Inference economics. `inference_ms`, `backend`, `node_id`, & sampling parameters aggregate into a routing intelligence layer. Our cache knows which nodes answer quickly, which backends run which models efficiently, & which quantization levels trade answer quality for speed. Cache hits carry this history forward — a future routing decision can prefer a node whose cached answers already proved useful.

Git-versioned knowledge graphs. For codebase queries, our `git_commit` field anchors every answer to an exact repository state. Our cache accumulates a version-annotated knowledge graph of our codebase: "at commit abc123, this function does X; at commit def456 it does Y." When a developer asks "what changed?", our cache surfaces the delta between two commit-keyed answer sets without re-reading any file.

9b. Unfirehose Integration (Proposed)

Unfirehose collects JSONL from Claude Code, Orchestra, uncloseai, & agnt.gg harnesses across our mesh. Our providence cache layer does not yet exist in production. This section describes our proposed integration.

We propose extending unfirehose with a `providence_cache` table & three new API routes:

`GET /api/providence?uri=...`

Returns all cached answer records for a document URI. Includes document root hash, question hashes, answer previews, & peer counts. Lets users audit what our mesh knows about a document.

`POST /api/providence`

Stores a new providence record. Accepts `{ document_root, document_uri, question_text, answer_text, merkle_proof, model }`. Validates our proof before storing.

`GET /api/providence/peers?root=...`

Queries our mesh for cached records matching a document root hash. Peer querying remains a future capability, opt-in per domain. Our initial implementation stays private: local cache only, no mesh broadcasting.

Our existing unfirehose dashboard would gain a Providence page: a view of our cache landscape across documents, cache hit rates per domain, & chain health.

10. Relationship to Zero-Knowledge Proofs

A Merkle proof constitutes a specific form of zero-knowledge argument: *I can prove this leaf exists in this tree without revealing our other leaves.*

In our context: *I can prove this answer came from this document version without revealing our full document.*

Our full ZKP apparatus (zk-SNARKs, Bulletproofs) provides stronger guarantees & more complex proofs. Merkle proofs provide a simpler, faster, & more practical subset: membership proofs sufficient for our provenance use case.

Our connection to ZKP deepens when our public chain handles sensitive documents. A user could prove their answer derived from a specific confidential document without revealing our document content, our question text, or our answer text. Only our proof path & our document root travel over our network. Our verifier confirms our proof without seeing our data.

This positions our public chain as opt-in verifiable knowledge sharing, not compelled disclosure.

11. Why This Matters for Permacomputer

Permacomputer infrastructure grows knowledge that outlasts our individual sessions, our individual machines, & our individual participants. Code seeds sprout on any abandoned technology. Knowledge should do our same.

A private chain that lives only on one machine dies with that machine. Our public chain on our unfirehose mesh survives node failures because our same document root on a different node still matches our cached proofs on our other nodes. Our knowledge persists as long as any mesh participant holds a copy.

Our Merkle structure aligns with permacomputer values:

Truth: proofs over assertions. Our answer carries a mathematical certificate, not just a model's confidence score.

Freedom: private chains for sensitive knowledge, public chains for shared knowledge. Our user chooses. No platform extracts our knowledge without consent.

Harmony: cache hits replace repeated inference. Our same compute produces more answers over time. Our mesh grows more efficient as our cache fills.

Love: small models with providence caches remove our infrastructure barrier. Hermes 3 Llama 3.1 8B on a consumer GPU, connected to our public chain, answers with our collective memory of our mesh. Our permacomputer does not require our most expensive hardware to deliver our most reliable answers.

12. Implementation Roadmap

Phase 1: Local private chain (unfirehose extension)

- Add `providence_cache` table to our unfirehose SQLite schema
- Extend our uncloseai.js Reverse RAG pipeline with pre-inference cache lookup & post-inference cache write
- Add `GET /api/providence` & `POST /api/providence` routes to our unfirehose web app
- Add Providence view to our dashboard

Phase 2: Public chain via unfirehose router

- Extend our router daemon with `/providence/broadcast` & `/providence/query` endpoints
- Add peer discovery to our mesh probe logic
- Add per-domain opt-in controls to our settings page
- Add peer contribution metrics to our Providence dashboard view

Phase 3: Proof verification in browser

- Implement `verifyMerkleProof()` in our uncloseai.js client
- Surface proof status in our chat UI (verified badge vs. fresh inference)
- Add document root display so users can audit our source fingerprint

Phase 4: Harness-level knowledge graph via unfirehose

Unfirehose collects JSONL from every harness session across our mesh. Every file read, every tool call, every assistant response produces a record. At the harness level, our Merkle Providence layer extends beyond web pages to entire codebases & research corpora:

- **Commit-keyed cache**: when a harness reads a file, we record `(git_commit_hash, file_path, question_hash)` $\rightarrow$ `answer`. Our same question about our

same file at our same commit hits our cache on every future session, on any node in our mesh, without re-reading the file.

- **Knowledge graph construction**: unfirehose builds a graph of concepts, files, functions, & sessions from harness JSONL. Nodes carry Merkle-verified summaries. Edges carry session provenance. Our graph grows with every harness run.
- **Graph-first traversal**: a harness navigating a large codebase queries our knowledge graph first. Only cache misses trigger fresh file reads & inference. Whether this reduces token usage, & by how much, remains for future measurement once a prototype exists.
- **Cross-session continuity**: a new harness session on our same codebase at our same commit inherits our full prior knowledge graph. No re-reading. No re-summarizing. Our mesh memory persists across sessions, machines, & agents.

This positions unfirehose not just as a log aggregator but as a shared, verifiable knowledge substrate for our entire mesh of machine learning agents.

13. Aborist: A Python Reference Implementation

Our uncloseai.js Reverse RAG client lives in our user's browser. Our Merkle Providence cache layer described above also belongs in our browser, & in a Next.js logger that aggregates JSONL across our mesh. A complementary surface exists at our other end of our pipeline: a server-side, content-addressed document store that ingests entire corpora at once, computes Merkle commitments at scale, & serves cached answers to any peer that produces matching cache keys.

Aborist (https://git.unturf.com/engineering/unturf/aborist) is our Python reference implementation of that surface. Our codebase ships under AGPL-3.0-only, runs against hermes.ai.unturf.com by default, & implements every primitive described in Sections 1 through 12 of this paper as working code with measurements on a real corpus.

13.1 Schema: v9.8 Admissibility Ledger

Aborist's providence cache uses an 8-dimensional cache key, a strict superset of our 6-dimensional key in Section 3. Two extra dimensions formalize what was previously implicit:

```
cache_key = source_root | question_hash | model_profile_hash
          | conversation_hash | governance_policy_hash
          | schema_version | canonicalization_version
          | chunking_version
```

Bumping any one stales every prior cache record, so a chunker change or a canonicalization rule change cleanly partitions our cache namespace. Records carry a `falsification_state ∈ {live, failed, stale, quarantined}`: queries always filter on `state='live'`, & a separate `falsifications` table logs every state transition with `cache_key`, reason, actor, & `audit_event_hash`.

Every state-changing operation (ingest, distill, falsify, evict, rehydrate, burn, mesh epoch rotation, snapshot creation) appends one row to a tamper-evident audit chain:

```
event_hash = sha256(prev_event_hash || canonical_json(body))
```

Verification is purely local: re-hash forward & confirm. A break in our chain signals tampering on our DB at a specific event index.

Two leaf-removal verbs cover different stages of corpus life:

- **Falsify** (audit-preserving). Flips `falsification_state` from `live` to `failed` / `stale` / `quarantined` and writes a row to `falsifications` (cache_key, reason, actor, timestamp, audit_event_hash). Cache lookups always filter on `state='live'` so the record stops being

reachable, but it stays on disk as forensic history. Use when downstream consumers may have ingested the answer.

- **Burn** (kindergarten leaf delete). Actually deletes the row. Kept off the default verb path because it forfeits history. Refuses if the leaf has children — for `providence_cache` leaves: any falsifications referencing the cache_key; for `documents`: any derivations using the doc as a source, any incoming edges, any providence records keyed on its document_root; for `cores`: same as documents plus downstream cores in derivations. `--force` overrides for cleanup edge cases & records `forced=True` plus the child count at burn time in the audit body. Always writes a `providence_burn` / `document_burn` / `core_burn` event regardless of forcing. Use during early/scratch corpus building before downstream consumers exist; otherwise prefer falsify.

Both surfaces (`aborist providence --falsify` & `aborist burn`) compose with cross-shard scope. Chain integrity after either operation is verifiable in O(1) per shard via `make chain-check-shards` — a `LEFT JOIN` query that counts dangling `prev_event_hash` references; 0 = chain intact, non-zero = a break to investigate.

13.2 Multi-Source ETL

Aborist ships with five source adapters, each yielding *Document* objects through our same *ingest_source* pipeline:

- **Wikipedia Phase III SQL dumps** (2003–2005). Hand-rolled streaming parser for MySQL extended INSERT syntax, bz2-aware, ~3 minutes per 128k articles on commodity hardware.
- **Wikipedia Phase IV XML dumps** (2006+). Streaming *iterparse* over *pages-articles.xml.bz2* with bounded memory; supports *pages-meta-history.xml.bz2* for full revision history. Tested against 2010-11 enwiki at 6.2 GB compressed.
- **xAI Grok account exports**. One Document per conversation, every message rendered with sender/model/timestamp inline.
- **Git & Mercurial repos** (self-play). Walks *git ls-tree HEAD* or *hg manifest*; one Document per text file. Re-ingest after new commits auto-chains via *supersedes* edges, so our Merkle tree grows alongside our repository.
- **HTML pages**. Robots-aware fetch with *selectolax* extraction; hyperlinks become outbound edges.
- **Live website crawl** (off by default, `aborist[crawler]` extras). BFS-discovers same-domain URLs from a seed via a verbatim lift of agents.ai.unturf.com's *AsyncWebFetcher* (ethical: respects robots.txt Disallow + Crawl-delay, with a *fast_mode* flag that drops the delay for trusted hosts). Discovered URLs are filtered: feeds (atom/rss/rdf/sitemap) and image/video/audio paths are crawl infrastructure, not knowledge content, & skip ingest. ETag and Last-Modified per response are captured into `document_http_meta` so a subsequent `aborist crawler recrawl-check` round can issue conditional `If-None-Match` / `If-Modified-Since` HEAD requests; 304 = fresh (no body transfer), 200 = stale (re-ingest), 404/410 = gone, other = unreachable. The crawl shard filename is derived from the seed hostname (`crawl_<host_with_underscores>.db` under `$(SHARDS_DIR)`) so cross-shard query picks up new domains without configuration.

Adding a seventh corpus is one new *Source* subclass. Our contract is intentionally minimal: yield Documents, our pipeline canonicalizes, chunks, hashes, persists.

13.3 Storage Cheats: 67% Reduction Measured

Three orthogonal optimizations stack against SQLite-resident corpora:

- **Zstandard-compressed chunk content.** Plaintext bodies compress 3 to 5x at level 3. Magic-byte detection on read makes legacy uncompressed rows transparent. Cores stay plaintext so SQL substring queries continue working.

- **WITHOUT ROWID on edges.** Our composite primary key covers every column, so a default rowid-keyed table near-doubles row data in our PK index. WITHOUT ROWID makes our table itself a B-tree keyed on our PK. We dropped our partial dst_uri index too, since only one analytic query actually filters on dst_uri.
- **Contentless FTS5.** Our prior schema had FTS5 store its own copy of every chunk's text. Contentless mode (*content=''*, *contentless_delete=1*) keeps only our inverted index. Snippets are built in Python by joining *chunks_fts.rowid* = *chunks.chunk_id* & decompressing. Requires SQLite 3.43+.

Apples-to-apples reingest of our 2003-05-16 enwiki cur dump (128,245 documents):

```
OLD schema:  2.6 GB across 4 shards    (~21 KB / doc)
NEW schema:  857 MB across 4 shards    (~6.7 KB / doc)
reduction:   -67%
```

Same document set, same Merkle roots, identical query results. Storage cost dropped by two-thirds with no change to our content-addressed identity.

13.4 Retrieval Quality: Recall Fixes

Our reverse-RAG client picks one document; aborist picks K from a corpus & merges. The ranking pipeline grew through five live-corpus catches; each fix earned its place by failing on a real query. The current shape is a multi-stage filter with three accept paths & a stem-tolerant token matcher.

13.4.1 Title-search baseline dominated FTS5

A single-token title overlap against a generic word like `"intel"` used to score 60+, swamping body-relevant FTS5 hits scoring 5-15. Result: every `Intel_*`-prefixed legacy article flooded our top-K for any "intel" query, even when a contemporary article like `Pentium_4` had body relevance & no title overlap. Our fix dropped our title-search baseline so title overlap competes with FTS5 instead of monopolizing it. Title contributes a token-overlap boost (`overlap × 10.0`) on top of BM25; it doesn't run as a separate ranking signal.

13.4.2 BM25 short-doc bias surfaced stubs over topical articles

A 1-paragraph stub mentioning every query token once outranked a 30-page topical article where our same tokens are diluted by length. We added a body-coverage rerank using `sum(sqrt(count(t))) * 0.6` over query tokens for each surviving candidate. Sqrt scaling differentiates stubs from topical articles without letting enumerative list pages run away with our score.

Verified on our 2010-11 enwiki corpus (3.47M documents, no distill cores):

```
before:  Intel_4040 (1974 microcontroller stub) ranked #1
         for "what is the fastest intel CPU?"
after:   Intel_Core_2 #1, Intel_Atom #2, Intel_4040 #3
         answer cites Core i7 specifications correctly
```

13.4.3 Three accept paths in the title-relevance filter

A hit clears `_filter_by_title_relevance` if any one of three accept paths fires:

1. **Title-token overlap.** Strongest signal. Direct stem-aware match between the query's title-tokens and the document's title-tokens. Synonym fallback (`concepts.synonym_expand`) only fires for 1-token queries — multi-token queries that admit synonyms over-recall (an "amd"-querying agent picking up *Intel_** titles via the Intel/AMD rivalry group).
2. **TF-IDF core keyword overlap.** Closes the gap for neologisms that never appear in titles but are distinctive enough to be the TF-IDF top-keywords of conversation bodies. `"permacomputer"` was the canonical case — never a Wikipedia title, but a strong TF-IDF signal in Grok-conversation cores. Without this path, neologism queries returned empty.

3. **Body density.** `_body_density_passes` admits docs that mention at least `breadth_threshold` distinct query tokens AND have `total_mentions >= 3`. Cheap proxy for "actually about the topic" when neither title nor core caught it.

Rivalry exclusion (`rivalry_excluded` in `aborist/qa/concepts.py`) applies on every accept path: an Intel-titled doc dropped from an AMD query, & vice versa. Caught the case of competitor docs leaking into single-side queries when both share a topic page.

13.4.4 Breadth-required filter — multi-token queries demand co-occurrence

Pre-2026-04-29 the body-density rule was "at least HALF the tokens present" — too lenient for 2-token queries. Fox caught this on 2026-04-29: `"who is supermans girlfriend?"` admitted seven unrelated `Girlfriends`-titled articles (TV show, films, songs) because they had `girlfriend` in body but no `superman` whatsoever; the half-rule passed at 1-of-2.

The fix scales the breadth threshold with query length:

```
≤ 2 query tokens   require ALL distinct tokens present (in body and in title)
3+ query tokens    require N - 1 (one weak signal token may miss)
```

Plus the depth check (`total_mentions >= 3`) still gates body-density. Title-token-overlap shares the same scaling. The "3+ require N-1" path lets one query token miss for long queries where models throw filler in (`"what is the plot of the matrix film?"` admits docs that have all of `plot`/`matrix`/`film` even when `what` is filtered out as a stopword & `the` doesn't count).

Verified on the same enwiki corpus:

```
before:  Girlfriends (1991 album), Girlfriend (Avril Lavigne song),
         Girlfriends (TV series), Girlfriends Films, ...
         — 7 of 8 top hits unrelated to Superman
after:   Lois_Lane #1, Superman_in_film #2, Smallville #3, ...
         — Superman/Lois Lane mythology surfaces correctly
```

13.4.5 Stem-aware token matching — possessive & plural collapse

Same `supermans girlfriend` defect had a second cause: `superman's` from the question canonicalized to `supermans` (apostrophe stripped by `_TITLE_TOKEN_RE`), which doesn't substring-match the bare `superman` in source bodies. Plurals had the symmetric problem: `girlfriends` query token vs `girlfriend` source mentions.

`_stem_token_for_match` is a one-line lite stemmer: trailing `s` strip on tokens of length > 4 that don't end in `ss`. Conservative — short tokens (`"is"`, `"as"`, `"us"`) keep their `s`; `"miss"`/`"class"`/`"chess"` are skipped. Both query tokens & document tokens pass through it before set comparison & body counting; matches in either form score.

This is a soft normalization step BEFORE the hard hash. `question_hash` itself does NOT stem — it canonicalizes by article-strip + trailing-punct-strip, which is enough for the equivalence-class layer. Stemming lives only in the retrieval ranker, where over-recall is preferable to under-recall (filtering happens later).

13.4.6 Per-source context cap — no single doc monopolizes the budget

Fox 2026-04-29: a query for `"who is batman?"` returned 8 sources but the LLM cited only 2 in its answer. Cause: `List_of_Batman_comics` (an 80 KB+ bibliography) was the top FTS5 hit & consumed the entire 60 KB `max_context_chars` at hit #1, leaving every subsequent source dropped with `char_budget <= 0`.

The fix caps each top-K hit at `max_context_chars / top_k` chars before respecting the global budget. With `top_k=8` and `max_context_chars=60000`, each source ships at most 7,500 chars. Bibliographies still appear in the source list but as one slice among many; the actual character article (`Batman`), the franchise page, the film article, the rogue's gallery, etc., all reach the LLM. Multi-source diversity restored without raising the global budget.

```
before:  8 sources retrieved, 1 in context (List_of_Batman_comics ate 60 KB)
after:   8 sources retrieved, 8 in context (each capped at 7,500 chars)
```

13.4.7 Worked example: dinosaurs in Jurassic Park

Query: `"what dinosaurs where in jurassic park film?"`. After stopword + stem: `{dinosaur, jurassic, park, film}`. 4 tokens, so `title_breadth = N - 1 = 3`.

```
title-overlap accept (3 of 4 stems in title):
    Jurassic Park (film)         {jurassic, park, film}
    Jurassic Park (film score)   {jurassic, park, film}
    Jurassic Park III (film score) {jurassic, park, film}

body-density accept (≥ 3 of 4 distinct tokens in body, total ≥ 3):
    Jurassic Park (franchise)
    List of Jurassic Park characters
    Jurassic Park: Operation Genesis
    Jurassic Park video games
    The Lost World: Jurassic Park
```

8-source context, each ≤ 7,500 chars after wikitext-base strip. The model produced 4 quoted spans (Brachiosaurus, Velociraptor, T-rex, Dilophosaurus); 3 verified verbatim, 1 contained the model's own `[...]` elision marker & failed substring match. Final classification: HYBRID / quote / 3-of-4. The unverified span surfaces in `aborist inspect`, where an operator can decide whether the elision is benign or worth re-asking under a tighter prompt.

13.4.8 Phrase-pattern route — closing the allusion / reference-frame gap

The three accept paths in 13.4.3 (title-token overlap, TF-IDF core, body density) all key on individual content tokens. They miss a class of query whose diagnostic signal lies in a verbatim multi-token sequence rather than in any single token. Caught live 2026-05-01 by fox on:

```
Q: "has oceania always been at war with east asia"
```

Pre-fix, the system surfaced literal-geography articles (`Oceania`, `Asia`, `Far East`, `East Germany`) & answered as if asked about real-world regions. The intended frame: George Orwell's *Nineteen Eighty-Four*, where the Party's propaganda treats Oceania as eternally at war with Eastasia despite alliances changing across the novel's continuity. The 1984 article exists in the 2003 enwiki corpus, but its title shares zero content tokens with the question, so even FTS5 phrase MATCH on `"always been at war"` recovered the article only to have `_filter_by_title_relevance` drop it before rerank.

Two paired changes close the gap:

1. **Phrase-pattern search route in** `_search_corpus`. For each n-gram extracted from the question (n=6 score 100, n=5 score 90), run an FTS5 quoted-phrase `MATCH` & add hits to the candidate pool. Stopwords stay IN the n-gram — diagnostic value of an allusion lives in the exact sequence including function words (`"always been at war"` >> `"always war"`). 4-grams got tried & rejected: `"always been at war"` matches generic war-history articles too noisily for downstream rerank to separate. 5+ tokens trade recall for precision; most allusions (`"may the force be with you"`, `"winter is coming"`, `"to be or not to be"`) survive at length 5 or higher.
2. **Accept-path 4 in** `_filter_by_title_relevance`. Phrase-route hits get a `phrase_match_roots` set threaded back from search; the filter passes any hit whose `document_root` lives in that set, regardless of title-token overlap. Without this path, the route from #1 retrieves the right article only to have it filtered out before it can rerank into top-K. Mirrors the existing TF-IDF core accept-path: a structural signal that bypasses surface-token gates.

A latent bug surfaced & got fixed alongside: `_search_corpus` previously returned a bare list while the caller did `getattr(hits, "_core_match_roots", set())` to fish out a sidecar set, but the sidecar attribute was never attached. The TF-IDF core accept-path silently received an empty set for an unknown duration. `_search_corpus` now returns a tuple

(hits, core_match_roots, phrase_match_roots, root_to_shard) so both routes thread correctly.

Verified live:

```
before:  HYBRID 2/16, sources: Oceania, Asia, Outline of Oceania, Far East,
         Decolonisation of Oceania, Central Asia, East Germany, ...
         — Nineteen Eighty-Four nowhere in top-K
after:   EVIDENCE-LINKED 1/1, citation [E13 | Nineteen Eighty-Four | 682f0a11]
         answer: "The text does not directly state... The passage describes
                  a change in alliances, where Oceania switched from being
                  allies with Eastasia to being allies with Eurasia, and
                  the public was manipulated to accept this change..."
```

The route generalizes beyond Orwell. Any verbatim 5+ token sequence from the question that appears in an article body lifts that article into consideration, regardless of whether the corpus document carries an allusion-bearing title. Per fox's 2026-05-01 review: "*A reference query often requires answering under multiple frames, not picking one frame and discarding the others.*" The retrieval side now surfaces the right frame; the answer-compilation side (multi-frame polarity contract) lives as Ticket #000002 in docs/ & remains future work.

13.4.9 Template-phrase stopwords on FTS5 MATCH

A second 2026-04-30 retrieval miss exposed the dual: "tell me all there is to know about york england?" returned Charles I of England, Elizabeth I of England, Church of England, etc. Pre-fix, AND-mode FTS5 expanded the question to "all" AND "there" AND "know" AND "york" AND "england", which favored "X of England" articles that incidentally mention every template token over the bare York article whose body lacks "all"/"there"/"know".

_FTS5_STOPWORDS (in aborist/search/fts5.py) & _TITLE_STOPWORDS (in aborist/qa/query.py) now strip tell show describe explain summarize say give list find make please all there know everything anything something plus the standard English stopword set. They must stay in sync — both filter the same set of question-shaping tokens. Retrieval-time only: stopwords don't enter cache_key or governance_policy_hash.

13.4.10 Title-LIKE search ORDER BY token-hits

A 2026-05-02 retrieval miss surfaced two compounding defects in the title-LIKE pass:

```
Q: "what date did back to the future come out?"
Pre-fix top-8: trilogy, ride, animated-series, The Future,
               Future Husband, The Day You Come, Ratchet & Clank
               Future, Future plc — film article entirely absent.
```

1. _search_titles had no ORDER BY. SQLite returned rows in arbitrary internal order; LIMIT 32 truncated before the genuine matches landed. The Back to the Future film article (title contains BOTH back AND future) lost the LIMIT race to substring-match junk: Out (poker) (single-token match on out), Aberdeen, South Dakota (south contains out), Backplane (back substring), Outline of biology (out substring), etc.
2. The caller's post-filter check used substring match (if t in title_lower) — same defect at a higher layer. "south" contains "out", "outline" contains "out", "backbone" contains "back" — all passed the no-op overlap check.

Two paired fixes:

a. The SQL builds a token-hit score via CASE WHEN LOWER(title) LIKE ? THEN 1 ELSE 0 END per token & ORDERs BY (title_score DESC, LENGTH(title) ASC). The film article's score = 2 (back+future) ranks above garbage's score = 1. Title length asc tie-breaks toward title purity (shorter title + same overlap = higher per-token signal).

b. The post-filter intersects token-sets stem-aware via `_title_query_tokens` & `_stem_token_for_match`. Word-boundary semantics; only genuine title tokens count.

Verified live: post-fix the BTTF film article ranks #1, the trilogy at #2, junk gone from top-8.

13.5 Federation: Mesh Layer (Off by Default)

Two peers that ingest our same dump compute bit-identical document_roots & cache_keys. Our mesh layer in aborist is our wire-and-trust scaffolding that lets them gossip those identities, dedup answers, & cleanly distrust an evicted member.

Cryptography (via *cryptography* library):

- **Ed25519** signs every membership mutation & every gossip envelope.
- **X25519 ECDH** wraps each epoch's symmetric mesh secret to every current member's DH pubkey. HKDF-derives a 32-byte AEAD key from each shared secret.
- **ChaCha20-Poly1305** authenticates payloads & per-member secret wraps; bodies optionally encrypt under the per-epoch shared secret with envelope metadata as AEAD additional data.

Membership state is per-epoch. Adding a member, kicking a member, or rotating our secret bumps our epoch & writes an audit event. Eviction works by rotating to a new epoch whose envelope omits our kicked member: their prior signatures stay verifiable forever (their old roster row is preserved), but they have no entry in our new envelope, so any AEAD-protected gossip from epoch+1 onward is opaque to them. Re-admission requires a fresh add operation by an admin.

Off by default: a *mesh.enabled* flag in our meta table gates everything. Fresh installs report *enabled: false* & no network code paths execute until our user runs *aborist mesh init* & *aborist mesh enable*.

13.5.1 Wire protocol

The HTTP gossip wire ships in `aborist/mesh/wire.py`. A peer runs `aborist mesh serve --host 0.0.0.0 --port 8400` to bind a stdlib `ThreadingHTTPServer`; another peer's client (`aborist mesh sync --peer http://...` and `aborist mesh pull --peer http://... --root <hex>`) speaks five message types, all wrapped in a single signed envelope:

```
ANNOUNCE_ROOT          document_root + source_uri + chunking/canonicalization/schema versions
ANNOUNCE_DERIVATION    core_root <- surface_roots(s), distiller_id, proof_blob hash
ANNOUNCE_PROVIDENCE    cache_key + audit_mode + answer_hash + verifier counts
ANNOUNCE_FALSIFICATION cache_key + reason + witness signature
REQUEST_BODY/DELIVER_BODY  body pull on local miss + per-chunk leaf hashes for Merkle re-derivation client-side
```

Receivers verify the Ed25519 signature against the sender's `sign_pub` looked up in `mesh_roster` at the envelope's `epoch_id`; non-members of that epoch produce no valid signature & their messages are silently dropped (HTTP 401, no audit event). Each accepted `ANNOUNCE_*` writes one `mesh_received` event into the local audit chain whose body is the full signed envelope, preserving non-repudiation. `REQUEST_BODY` triggers a `DELIVER_BODY` carrying the document text + per-chunk leaf hashes; the client re-derives the Merkle root from those leaves & refuses to insert any body whose leaves don't reconstruct the requested root.

13.5.2 Chain-of-claims tracking & fork detection

Per the protocol contract, audit chains across peers cannot "merge" in the literal sense — each peer's chain is a strict-order log of its own actions — but a receiver can detect when a sender broadcasts inconsistent histories. Each gossip envelope carries:

- `prev_event_hash` — the SENDER's previous event_hash they broadcast (or null for the first message);
- `event_hash = sha256(prev_event_hash || canonical_envelope_bytes)` — the sender's deterministic chain-of-claims hash for THIS message.

The receiver's local `mesh_peer_chains` table tracks, per peer, their last-known `event_hash` & sequence number. On accept: verify Ed25519 signature → recompute `event_hash` from canonical bytes & assert equality with the claimed value (400 on mismatch) → look up the sender's prior `event_hash` & assert equality with the claimed `prev_event_hash` (409 Conflict, "fork detected", on mismatch) → write the `mesh_received` audit event & advance the per-peer chain row. The first envelope a receiver sees from any peer is accepted with any `prev_event_hash` value because no prior is on file.

Catchup (peer sends multiple events to fill a gap) is deferred to a future commit; v1 raises a 409 on missed events & relies on the operator/agent to retry. The wire-level fork detection composes with the local audit chain: a tampered DB locally fails the `event_hash` re-hash chain at ingest time, & a forked sender fails the `prev_event_hash` chain at the first received message after the fork.

13.5.3 Optional AEAD body encryption

Signing the envelope guarantees integrity; the body is cleartext by default for backward compatibility. Senders can opt body content into AEAD encryption under the per-epoch shared secret: at construction time, body content is canonicalized to JSON bytes, ChaCha20-Poly1305 encrypted with a fresh 12-byte nonce & `f"{type}|{sender_id}|{epoch_id}|{ts}"` as additional authenticated data, & the resulting `{nonce_b64, ct_b64}` lands in a new `encrypted_body` envelope field while `body` becomes a placeholder. Both fields participate in `canonical_bytes()`, so the Ed25519 signature binds the ciphertext — a MITM cannot swap encrypted_body without invalidating the envelope.

Receivers detect encryption by the presence of `encrypted_body`, recover the epoch shared secret via the existing `mesh.state.recover_epoch_secret` (audit-chain → rotator → ECDH-unwrap-own-slot, no new sender-side cache needed), AEAD-decrypt, & continue normal handler logic with the recovered cleartext body. A receiver missing the epoch secret (evicted member, never enrolled) returns 401 — confidentiality failure is indistinguishable from auth failure at this layer. Mixed-mode peers interoperate: cleartext envelopes still verify byte-identically to the v1 wire format because `encrypted_body=null` is dropped from the canonical form before signing.

13.5.4 What converges across peers

Three classes of derivative reconcile across peers:

- **Deterministic distillers** (e.g., first-sentence cores): same src_root → same core_root, on every machine. Shareable as-is via the existing *derivations.proof_blob*.
- **Corpus-statistical distillers** (e.g., TF-IDF keywords): convergent only under shared corpus scope. Requires a *scope_root* commitment so peers explicitly agree on which corpus statistics underlie our cores.
- **LLM-derived answers**: non-convergent (serving non-determinism), but cryptographically witnessable. Same cache_key from many peers + the same *answer_hash = sha256(answer_text)* collapses to one row with many witnesses; divergent answers stay as multi-witness candidates, each Merkle-bound to a verifiable context_root.

The full operator runbook for a two-host federation lives in `docs/mesh-deploy.md` & the protocol contract (with diagrams) in `docs/mesh.md` in the aborist repo.

13.6 Corpus Snapshots: One Hash Names Our Forest

A *snapshot_root* is *MerkleTree.build([sorted DISTINCT document_roots]).root*: one 32-byte hash naming our content-addressed forest at a point in time. Two peers that ingested our same dump compute bit-identical snapshot_roots, so cross-machine "are we synced?" becomes a single hash comparison instead of doc-by-doc gossip.

Snapshot creation pins our root into our audit chain & records *(snapshot_root, taken_at, audit_event_hash, doc_count, parent_snapshot, reason)*. Snapshots auto-link to our most recent prior snapshot, giving a chain for free. The same *snapshot_root* doubles as our TF-IDF *scope_root*, so derivative reconciliation in 13.5 maps directly onto snapshot agreement.

13.7 Cross-Machine Convergence: A Worked Example

Our 2010-11 enwiki corpus, ingested on a single workstation in 1h 41m (4-way sharded), produces:

```
snapshot_root = 43797e46605de08dbab06cdcaf5be7ad78243b193c56f8580200dee6bcc7e1b9
doc_count     = 3,468,134 (deduplicated)
storage       = 38 GB across 4 shards
chunks        = 6,235,672
edges         = 90,593,523
audit_events  = 3,468,308 (chain intact, 30/30 random Merkle proofs verified)
```

A query against this corpus for "tell me plot of alien film" returned a cache hit:

```
document_root = 896c27f8d9de8e0fc8c9b337274e794bea4206dccd50b9654b34833d9d29f005
document_uri  = https://en.wikipedia.org/wiki/Alien_(film)
context_root  = 896c27f8...f005   (single source, equals document_root)
cache_key     = 1aa20457e980f9e564641b7c3a86edbc90bd9fb31d6e461e9da2fdc4082ddda1
answer        = (plot summary verbatim from the article)
```

Every one of those four hashes is bit-identical on any machine that ingests our same *enwiki-20101011-pages-articles.xml.bz2* (verifiable via our *md5sums.txt*) with our same aborist commit. Hand *896c27f8…f005* to any peer; if their corpus matches ours, they compute our same hash, with mathematical certainty. No trust required, just shared inputs.

What does NOT converge across two machines:

- *audit_events.event_hash* chain (event order is per-instance witness, not content)
- *documents.ingest_ts* & *chunks.chunk_id* (per-instance metadata)
- *answer_text* under fresh LLM inference (Hermes serving carries tiny non-determinism)

Everything in our cache key dimensions converges. Everything that observes our corpus's processing diverges. That separation is our admissibility property: one peer's witness of "what was in our forest at time T" is bit-equal to every other peer's witness of our same forest, while one peer's audit chain proves *that peer* witnessed our forest at our times *they* recorded.

13.8 Audit Mode: Source vs. Emergence

A Merkle proof binds an answer's *context* to a content-addressed root. It does not bind the *answer* to that context. A model handed a chunk of episode plot summaries can correctly recite the entity names that appear in the chunk, then complete the response with character bios drawn from training data — names that match, claims that emerge. Both layers are real. Conflating them under one "STRICT" label overclaims the proof.

v9.8 names the trichotomy (Section 13.1 inherits it from Merkle-AGI's audit contract): STRICT, HYBRID, VISUAL. Aborist decides which label applies *per answer*, by post-LLM faithfulness check, never as a default. The substrate's "VISUAL" label was named for neural-network proof systems where it meant "FOR-style visualization, no formal guarantee attached." In a RAG layer that semantic reads as *no recoverable proof of grounding*, so aborist persists the label as **UNGROUNDED**:

```
STRICT      every evidence unit in answer verifies verbatim against context
HYBRID      some claims source-grounded, some emerged from training
UNGROUNDED  no evidence, or none verifies — purely emergent
```

13.8.1 Layered verifier

aborist/qa/verify.py runs four strategies in sequence. The first to find evidence classifies the answer; later strategies don't run. Each strategy is a lexical operation under norm-v1 canonicalization (NFC, collapsed whitespace) plus case-folding — substring tests for strategies 1-3, token-coverage for strategy 4. The soft-hash boundary from CLAUDE.md applies: embeddings train rankers, hashes carry proofs, & we never mix the two. `verifier_method` on the providence record names which path fired so the audit chain stays diagnostic.

Strategy 1: quote. The system prompt instructs the model to wrap every factual claim in a verbatim double-quoted span from a source. The verifier locates every double-quote character in the answer & pairs them sequentially: 1st & 2nd, 3rd & 4th, & so on. Each pair brackets one quoted span; text between consecutive pairs is the model's framing prose & is not captured. Pure-regex pairing fails on adjacent quote pairs like `"title" prose "quote"` — the regex captures `prose` as a phantom span because every `"` looks like both an opener & a closer to it. Sequential pairing eliminates the artifact.

Strategy 2: span. When the model declines to quote but writes lines verbatim from source, bullet & sentence units from the answer are substring-tested against context. Catches paraphrased-into-prose answers where every claim is grounded but the formatting drifted from explicit quote marks.

Strategy 3: entity. When neither quote nor span lands evidence, multi-word capitalized phrases (Carrie-Anne Moss, Thomas A. Anderson, Agent Smith) are extracted from the answer & substring-tested. Catches the Wikipedia-infobox case: source has structured `[[Keanu Reeves]] [[Laurence Fishburne]]` markup; model paraphrases into prose; spans diverge but every named entity is intact & grounded. Entity-existence in source is, however, weaker than claim-level evidence — the model could correctly name entities while making *claims around them* that came from training. Aborist gates the entity path with a policy:

```
strict     all entities verify → STRICT (legacy; overclaims TMNT-style cases)
hybrid     any entity verifies → HYBRID (honest cap)
drop       skip entity path entirely → UNGROUNDED (most conservative)
proximity  STRICT only if N=3 verified entities cluster within W=300 chars
           in source. Otherwise demoted to HYBRID/UNGROUNDED.
```

The default policy is **proximity**. Entity-cluster density distinguishes "source documents these entities as a group" (cast list, infobox, roster) from "source mentions them incidentally in scattered prose." A Wikipedia infobox passes the cluster test & promotes; a TMNT episode plot summary that names the four turtles & Splinter across separate sentences fails it & demotes. The policy lives in `DEFAULT_QUERY_POLICY["entity_policy"]` so any change folds into `governance_policy_hash` & invalidates prior cache cleanly.

Strategy 4: paraphrase. The span strategy's secondary path. When a sentence-shaped span fails substring match in strategy 2, a token-coverage probe runs on the same span: tokens of length ≥ 4 chars are extracted (less per-token punctuation strip so `wayne,` lines up with bare `wayne`), filtered against an English stopword set (`from`, `with`, `have`, `would`, `which`, `where`, `their`, etc. — high-frequency fillers that match almost any English text), and the fraction of remaining content tokens present anywhere in the normalized base context is computed. The span is accepted as paraphrase-verified when coverage clears `DEFAULT_PARAPHRASE_COVERAGE` (default 0.85) with at least `DEFAULT_PARAPHRASE_MIN_TOKENS` (default 4) content tokens. The label `verifier_method='paraphrase'` distinguishes records where any soft-verified items contributed from records where every match was lexical-verbatim. Paraphrase fires only on prose-shaped spans, not on lists of proper nouns — a span needs at least 2 lowercase content tokens (`_is_prose_span`) to qualify, so a list like `"Keanu Reeves, Laurence Fishburne, Carrie-Anne Moss"` continues through to the entity strategy where proximity policy can tight-cluster check.

The stopword filter is calibration choice, not threshold relaxation. Including stopwords in the denominator inflates coverage scores for cases where topical content is missing — a fabrication like `"Batman is the alias of Bruce Wayne, a wealthy businessman who resides in Gotham City"` (where `wealthy` / `businessman` / `resides` aren't in source) would score ~0.67 raw but should NOT promote. Filtering the stopwords drops the denominator to topical tokens only; the missing topical tokens dominate the score & the span correctly stays UNGROUNDED. Lowering the threshold itself (e.g. 0.85 → 0.67) would have promoted the fabrication; tightening the signal preserves discrimination.

Token-coverage is a soft signal: it answers "the same content tokens appear in source," not "the answer is a byte-equivalent transcription." The hard chain still records the trichotomy classification, but auditors reading STRICT can use `verifier_method` to tell verbatim-from-source-text from same-tokens-different-sequence. Quote strategy (strategy 1) deliberately does NOT receive paraphrase fallback — the double-quote convention asserts verbatim citation, & paraphrasing-then-quoting is a model error to flag, not auto-promote.

13.8.2 Wikitext base prose

Aborist stores raw MediaWiki wikitext in chunk content so the link graph & original markup are recoverable from any page. The LLM produces clean prose. Without normalization, every wikilink-carrying source paragraph compares as a different surface form & the verifier wrongly reports UNGROUNDED on genuine source-grounded quotes — the FF7-characters page was the input that triggered the design (six legitimate quotes flagged as ungrounded because the model wrote `Cloud Strife` & the source had `[[Cloud Strife]]`).

`aborist/wikitext.py` provides `to_base(raw)`, a pure function pinned by `BASE_VERSION` (`wikitext-base-v1`). Same wikitext → same prose, forever, until `BASE_VERSION` is bumped. The algorithm runs `mwparserfromhell`: drop `<ref>` tags & namespace-prefixed wikilinks (`[[File:...]]`, `[[Image:...]]`, `[[Category:...]]`), call `strip_code(normalize=True, collapse=True)` to reduce surviving markup to display text, then collapse whitespace runs. Optional dependency — installations without `mwparserfromhell` fall back to identity & both LLM and verifier proceed against raw wikitext (the strip becomes a no-op rather than a hard requirement).

The conversion runs in **two places**, gated by the same `policy["base_version"]` flag:

1. **Before the LLM call.** `aborist/qa/runner.py` and `aborist/qa/query.py` apply `to_base()` to the assembled context before the user-turn arrives at Hermes. The model never sees `[[wikilinks]]`, `{{templates}}`, or `<ref>` tags — it sees the prose a human reader would see. Two consequences: the model can quote source paragraphs verbatim instead of escaping markup it doesn't understand, & Wikipedia chunks ship with ~43% fewer tokens (measured against the FF7-characters chunk). The token reduction compounds across context windows packed with multiple sources.
2. **Inside the verifier.** `verify_quotes` re-applies `to_base()` to the context before substring-testing the model's quoted spans. Idempotent — passing already-stripped prose through `to_base()` returns it unchanged. Same canonicalization on both sides means the verifier compares like-against-like.

`policy["base_version"]` lives in `DEFAULT_POLICY` and `DEFAULT_QUERY_POLICY` so it folds into `governance_policy_hash`. Bumping `BASE_VERSION` (e.g. `wikitext-base-v1` → `v2` after an algorithm change) invalidates every prior cache record's 8-dim key on next lookup; no schema migration, the next `ask` re-derives against fresh prose.

13.8.3 Emergence is preserved, not suppressed

A HYBRID or UNGROUNDED record stays live in cache. Training-derived knowledge is often correct (the turtles really do wield katanas), just unverifiable against the current corpus. We label honestly so a downstream consumer can route: STRICT answers carry compliance-grade provenance, UNGROUNDED answers carry *the model believes this, no certificate attached*. A model that genuinely refuses ("I don't know based on the provided sources") classifies UNGROUNDED with `verifier_method='none'` & `n_quotes=0` — the most honest answer the system can produce.

Emergence is a corpus-growth signal. Unverified spans persist on each providence record; `aborist emergent --aggregate` ranks them by frequency. Spans the model produces repeatedly that the corpus cannot ground are candidate ingest targets — what training has that ingestion is missing. When the missing source eventually arrives, `aborist reclassify` re-runs the verifier against existing answers without any LLM call, persisting the upgraded label & writing one `providence_reclassify` audit event per changed row. HYBRID promotes to STRICT, UNGROUNDED to HYBRID, & the audit chain remembers the transition. Prior records stay as historical witnesses: *on this date, against that corpus root, the model emerged this much beyond its sources.*

13.8.4 Verifier stays binary; falsifications carry soft signal

The verifier returns evidence units & a classification — that's the entire output surface. No per-quote diagnosis fields, no match-confidence scores, no partial-match indicators on hard verifier output. "Why didn't this ground?" lives on the falsify+reclassify loop where an operator (or an automated process) inspects the unverified spans, decides whether the failure is a corpus gap or a model fabrication, & either ingests new source or quarantines the record. Keeping the binary discipline on the verifier preserves the soft-hash boundary: classification is a hard bit, not a soft score.

The SOURCE/EMERGENCE distinction also clarifies the v9.8 cache_key invariant. Bumping `governance_policy_hash` (e.g. by tightening the system prompt or flipping `entity_policy`) invalidates cache lookups but leaves the source merkle tree, chunks, document_roots, audit chain, & snapshots bit-identical. Q&A records under the old policy stop being reachable; corpus identity is unchanged. No re-ingest required — the 8-dim key partitions cleanly along architectural seams.

13.8.5 Trailing-citation strip

Models trained on cited text often append parentheticals like `(Source: https://...)` or `(citing Wikipedia)` after a verbatim source sentence. Without preprocessing, the verifier substring-tests the prose **plus** the appended citation; the citation isn't in the corpus, so the whole span fails & the answer flags HYBRID or UNGROUNDED. The prose itself was verbatim.

`extract_quotes` and `extract_claim_spans` strip a single trailing parenthetical at end-of-span before substring testing, gated on a conservative regex: the parenthetical contents must start with a citation cue word (`Source:`, `citing`, `see`, `ref`, `reference`, `from`) OR contain a URL. A genuine prose parenthetical like `Pikachu (a Pokémon species) lives in forests` carries no cue word & no URL & stays untouched. The strip recovers verbatim-source spans without weakening the verifier on real disagreements.

13.8.6 Sidecar diagnostics

The verifier's output surface is deliberately small: an audit_mode label, an n_verified count, the unverified spans, & the verifier_method that fired. No per-quote diagnosis fields, no match-confidence scores, no nearest-neighbor traces. Keeping the binary discipline preserves the soft-hash boundary — classification is a hard bit, not a soft score, & enters `governance_policy_hash` cleanly.

But operators still need to ask "*why* didn't this span ground?" when triaging the unverified list. v9.8's answer is a **sidecar**: a read-only diagnostic verb that pulls the same source chunks the verifier saw, runs the same canonicalizations, & classifies each unverified span against the corpus. Sidecars never write to `providence_cache`, never extend `audit_events`, never alter any v9.8 field. They run on demand, observe state, & report. The classification-as-hard-bit invariant survives because sidecars participate in no proof path — they help an operator triage; they do not change what's signed.

The reference sidecar in aborist is `aborist inspect --cache-key X`. Each unverified span lands in one of:

- `verbatim_in_base` — the span IS in the base context. Surfaces a verifier or canonicalization regression so an operator can fix the root cause.
- `verbatim_in_raw_only` — span matches raw wikitext but not the base form. Wikitext-strip edge case worth its own label.
- `trailing_artifact` — a long prefix (≥ 60 chars) matches; the tail is reported separately. Catches model-appended citations that the strip in 13.8.5 didn't reach (different cue words, multi-line, etc.).
- `paraphrase` — high token coverage (≥ 0.85) but not in this sequence. Model rewrote source content; ingest more sources or tune the prompt.
- `partial_paraphrase` — 40-85% token coverage. Mixed sourced/emergent content within a single span.
- `no_overlap` — < 40% coverage. Likely full invention; corpus-growth signal.

Two additional sidecar diagnostics ship alongside the unverified-span classifier; both share the same write-nothing discipline:

- `diagnose_deflection(question, answer)` — surfaces topic-shift on adversarial-premise questions. Empirically caught 2026-04-30 on `"who is a benevolent dictator for life for mars?"`: the verifier returned STRICT with an answer about Guido van Rossum & Python, never naming Mars. The answer grounded honestly in a real BDFL article; the user's actual question went unanswered. **Subject-anchor heuristic**: the LAST content token in the question (after stopword strip) gets treated as the question's primary subject. If the subject anchor doesn't appear in the answer, classify as `deflection` regardless of generic-vocabulary overlap (the BDFL case has 3/4 generic-vocab match — `benevolent`, `dictator`, `life` — & still misses `mars`). Suppressed for shape-mismatch cases (`when`, `what year`, `how many`, `why`) where the answer structurally need not echo the question's subject (`"1989"` answers `"what year did the berlin wall fall?"` correctly without re-naming `"wall"`). Wired into `inspect_cache_key()` & per-row in `bench/qa_sweep.py`; never feeds back into providence.
- `diagnose_title_relevance(claim, cited_source_titles)` — flags retrieval-driven hallucinations where the cited chunk's source title shares zero stems with the claim's content tokens. Empirically caught 2026-05-01 on `"explain spin glass modeling, tensors?"`: the verifier returned STRICT 1/1 on a claim about spin-glass modeling cited to a chunk from the *Quantum chromodynamics* article. The token-coverage check inside the chunk passed accidentally on shared physics vocabulary (`spin`, `model`); the article isn't about spin glasses. The 2003 corpus's *Spin glass* article (which exists) carries zero occurrences of `"modeling"` or `"tensors"`, so AND-mode FTS5 returned no hits & OR-fallback ranked unrelated articles by raw BM25. The sidecar surfaces the mismatch (`claim_tokens = {spin, glass, modeling, tensor, ...}`, `title_tokens = {quantum, chromodynamics}`, `overlap = ∅`) without modifying the binary verifier output. Promotion to a hard verifier check stays deferred pending bench evidence on false-positive risk for legitimate cross-document grounding.

The sidecar pattern generalizes beyond *inspect*. Any future tool that wants to surface *why-not-grounded* — embedding-based nearest neighbors, LLM-generated rationales, fuzzy-match scores — belongs in a sidecar. The hard chain stays binary; the soft channel proliferates as the project learns. Sidecars compose cleanly because they all read from the same v9.8 record & contribute nothing back.

13.9 Claim-Lattice-Pointer Mode (CTI / Quote-by-Pointer)

The quote/span/entity/paraphrase verifier in 13.8 audits answers the model wrote in its own words. It catches one defect class cleanly — *did the quoted text appear verbatim in source* — but it cannot catch a related class: **synthetic elision**. A model emitting `"prefix [...] suffix"` produces a span whose halves both verify on their own, glued together by a model-typed ellipsis that flattens our ground truth. The verifier reports STRICT on a frankenquote.

Aborist's response is a second answer mode that makes synthetic elision **structurally impossible**: the model never types the quote string. Instead, the runtime owns extraction, the model owns selection, & the renderer owns display. We call this *Clause Lattice Intelligence* (CTI), or in implementation terms `policy["answer_mode"] = "claim_lattice_pointer"`:

```
raw model output
→ clause extraction by pointer parser
→ claim nodes with evidence pointers
→ deterministic verifier
→ admissibility decision
→ rendered prose with literal source spans
```

The model writes natural prose with bracketed evidence-id tags:

```
Steve Jobs co-founded Apple. [E1]
Steve Wozniak co-founded Apple. [E1,E2]
Ronald Wayne co-founded Apple. [E1]
```

Each `[E\d+]` references an entry in a runtime-built **evidence map** the model saw in the prompt. The runtime carries two ids per evidence object — a short `pointer_id` (`E1`, `E2`, …, sequential within the run) shown to

the model, & a `evidence_id` (sha256-derived `E########`, content-addressed) used by the cache, run-DAG, & audit chain:

```
EvidenceObject {
    pointer_id    "E1"                                    # model-facing
    evidence_id   "E1f8e4c2a"                              # run-stable
    source_root   <document_root>
    chunk_root    <chunk leaf hash>
    offset_start  0
    offset_end    <len(span)>
    source_role   primary_answer_source
    text_hash     sha256(span)
    span          <literal source text>
}
```

Two-layer id discipline is load-bearing. Showing the model the hex form (`E1f8e4c2a`) tokenizes to 4-6 BPE tokens of out-of-distribution noise that nudges Hermes 3 toward DSL/code-completion mode. Showing it `E1` is one BPE token & lands inside the citation-style prose distribution that academic papers, Wikipedia articles, & footnoted text taught the model. Citation style is heavily represented in training; random hex is not. The runtime maps pointer_id → evidence_id at verification time; everything downstream of the verifier sees the content-addressed form.

13.9.1 Pointer lines vs JSON — both first-class, different defaults

JSON is one encoding for *(claim_text, [evidence_ids])* pairs but not the right *default* for an 8B instruction-tuned model. JSON discipline failures (missing braces, prose preamble before the object, escaping errors) generate spurious one-shot SCHEMA_INVALID verdicts even when the model knew the right answer. Pointer-line stays inside the model's prose distribution: the model writes prose, sometimes with bracketed citations, exactly as it has seen citations written ten thousand times. The parser walks lines, pulls every `[E\d+]` & `[E\d+,E\d+,…]` bracket payload, & returns `(claim_text, pointer_ids)` per line. Lines without a tag get `parse_status = "NO_EVIDENCE_POINTER"` & count toward the denominator so the model can't smuggle unsourced prose past the verifier. Lines with a tag but no surviving claim text after the strip get `parse_status = "SCHEMA_INVALID"`.

Both modalities stay first-class in the substrate. `answer_mode = "claim_lattice"` is the JSON variant; `"claim_lattice_pointer"` is the prose-with-citations variant. They produce equivalent claim-evidence lattices, hash to distinct `governance_policy_hash` silos, and produce comparable run-DAGs (the JSON variant skips the `parsed_claim_lattice` stage because the model's output IS the parsed form). The substrate exposes both so an agent can route its query to the modality that fits its inference path: pointer-line suits 8B prose-distribution models, JSON suits grammar-constrained inference (vLLM `guided_json`, Claude / GPT-4 native JSON mode) where the model emits valid objects deterministically. Aborist provides a lenient JSON pre-parser (strip markdown fence wrappers, trim non-object preamble/suffix, normalize curly quotes, fix trailing commas) so the JSON variant survives Hermes-class drift long enough for an agent to observe its SCHEMA_INVALID rate and decide whether to keep using it. Picking between modalities is an agent-level concern; the substrate's job is to expose them, telemetry both, and never silently pick winners.

13.9.2 Hard verifier — eight deterministic checks

`verify_claim_lattice` runs **only** deterministic checks. The hard/soft boundary from 13.8.4 holds — semantic entailment (NLI-grade), completeness, predicate compatibility, scope ambiguity, & counterevidence retrieval all stay sidecar territory. The hard checks:

```
1. parser_succeeded            line had a recognized [E\d+] bracket
2. evidence_id_resolves        pointer maps to an entry in the
                               runtime-built evidence map
3. source_role_allowed         resolved entry's source_role is in
                               the policy allowlist
4. claim_text_non_empty        the line carried prose around the tag
5. citation_coverage           claim's content tokens textually
                               overlap the cited span at coverage
                               ≥ min_citation_coverage (default 30%);
                               short claims (≤3 content tokens)
```

```
                                 require ≥1 token match. Catches the
                                 lazy-anchor failure where a claim
                                 cites a chunk whose text shares only
                                 a generic topical word with the claim
6. pointer_count_within_cap      len(pointer_ids) ≤ max_pointers_per_claim
                                 (default 2). Over-cited claims are
                                 trimmed to the first N pointers and
                                 verification proceeds; a
                                 POINTER_OVERFLOW_TRIMMED violation
                                 is recorded so STRICT is no longer
                                 reachable for the run. Trim-and-verify
                                 protects correct-but-over-cited claims
                                 without letting the over-citation
                                 pattern pass silently
7. anchor_class_warrant         five anchor classes compose. Each
                                 gates on either question shape, claim
                                 content, or both. A claim flunks
                                 warrant when ANY required anchor in
                                 ANY active class is missing from
                                 every cited span.
                                 (a) proper-noun anchors — gated on
                                     relation-question shape; AT LEAST
                                     ONE Title-Case anchor must appear
                                     in some cited span. Catches the
                                     relation lazy-anchor class.
                                 (b) date anchors — always-on, gated
                                     on the CLAIM containing a 4-digit
                                     year (1500-2199) or full English
                                     month name; ALL components must
                                     appear in some cited span.
                                     Catches the date lazy-anchor
                                     class.
                                 (c) entity-list anchors — gated on
                                     entity-list-shape questions
                                     ("name the X", "list the Y");
                                     AT LEAST ONE named entity from
                                     the claim must appear in some
                                     cited span. Demote-don't-reject:
                                     an extra entity from training
                                     prior is acceptable as long as
                                     cited evidence anchors at least
                                     one of the named entities.
                                 (d) count anchors — gated on
                                     count-shape questions ("how many
                                     X"); ALL count tokens in the
                                     claim must appear in some cited
                                     span in either word OR digit
                                     form (digit↔word equivalence:
                                     "two" ↔ "2", "twelve" ↔ "12").
                                 (e) cause anchors — gated on
                                     why-shape questions ("why did X
                                     happen"); AT LEAST ONE cause
                                     anchor (proper noun OR ≥5-char
                                     lowercase common noun outside
                                     a small stopword pool) must
                                     appear in some cited span.
                                 Vacuous-pass when no class fires —
                                 broad descriptive claims flow
                                 through unchanged.
8. title_relevance              cited evidence's source title must
                                 share ≥1 stemmed content token with
                                 the claim text. Per-claim, ANY-match
                                 across cited sources: the rule fires
                                 only when ALL cited titles miss.
                                 Demote-to-HYBRID via TITLE_MISMATCH
                                 violation. Catches the
                                 retrieval-driven hallucination class
                                 (2026-05-02 spin-glass case): claim
                                 about spin glass modeling cited to a
                                 chunk from *Quantum chromodynamics*
                                 passed Rules 1-7 on incidental
```

```
                                    physics-vocab overlap, but the
                                    SOURCE was never about the claim's
                                    subject. Vacuous-pass when claim or
                                    title has no extractable tokens.
```

Aggregation matches 13.8's trichotomy: STRICT = ≥1 verified pair AND zero violations of any kind; HYBRID = some pairs verified, some failed (or any soft-demote violation: POINTER_OVERFLOW_TRIMMED, LAZY_ANCHOR_DEMOTED, WARRANT_MISSING); UNGROUNDED = no verified pairs. Per-claim statuses sharpen the diagnosis without breaking the binary discipline: `EVIDENCE_LINKED`, `EVIDENCE_LINKED_PARTIAL`, `UNKNOWN_EVIDENCE_ID`, `SOURCE_ROLE_BLOCKED`, `CITATION_MISMATCH`, `NO_EVIDENCE_POINTER`, `SCHEMA_INVALID`.

Warrant-lite is the lattice's first **semantic-grounded** hard check, and it earns the proof-path entry by staying lexical. Regex shape-detectors decide which classes activate (relation, entity-list, count, why-cause); Title-Case proper-noun extraction, 4-digit-year extraction (1500-2199, skipping ZIP codes / elevations), full-month-name extraction, count-token extraction with digit↔word equivalence, and lowercase-common-noun extraction (≥5 chars, outside a small stopword pool) pull the candidate anchors; case-insensitive substring tests ask whether the required anchors appear in cited spans. No NLI, no embeddings, no LLM-as-judge. The check is reproducible byte-for-byte across runs and folds cleanly into `verifier_policy_hash` via `policy["claim_lattice_warrant_check_enabled"]`. The principle: **pointer verification is not warrant verification**. A pointer resolves to a span; a warrant requires the span to actually contain the claim's load-bearing anchors.

Warrant-lite generalized in May 2026 (Ticket #000003) to cover three additional question shapes whose lazy-anchor signature the original two classes missed:

- **Entity-list shape** (`name X`, `list X`, `who are the members of X`). List-aware extractor returns multi-word phrases ∪ solo-cap individual names so comma-separated entities each contribute to the anchor pool. ANY-match semantics (demote-don't-reject): an extra entity from training-prior is acceptable as long as the cited evidence anchors at least one named entity.
- **Count shape** (`how many X`, `how much X`). Digit ↔ word equivalence with ordinal collapse: a claim saying `"six"` matches a span containing `"6"`, & `"sixth"` collapses to cardinal `"6"` for the substring test. Year-shaped digits filter out (those belong to the existing date-anchor class). ALL-match semantics: every count token in the claim must appear in some cited span as digit or word.
- **Why-cause shape** (`why X`). Cause-anchor pool widens to ≥5-char lowercase common nouns post a quantifier-adjective stopword set (`various` / `factors` / `situation` filter out) plus the existing proper-noun anchors. Catches lowercase common nouns the relation-shape extractor misses (`iceberg` / `asteroid` / `propaganda`). Gated on why-shape only: lowercase common-noun extraction has higher false-positive risk elsewhere.

Each shape's anchors get extracted when the question shape matches & failures accumulate across classes — a relation-shape claim with both a year & a proper-noun must satisfy both. Per-class policy gates were intentionally not added: the single `warrant_check_enabled: bool` flag is the minimum-viable mechanism. Per-class flags earn their slot only when bench evidence shows over-firing on a specific class — see the five-step algorithm below.

Five 2026-05-01 failure classes motivated the design:

- *Relation lazy-anchor.* "Who is Homer Simpson's boss?" answered "Homer Simpson's boss is Mr. Burns. [E1]" cited to a voice-actor bio paragraph that mentioned "Homer" but never "Mr. Burns" — the structural checks all passed, the proper-noun warrant catches it.
- *Date lazy-anchor.* "What date did Back to the Future come out?" answered "Back to the Future was released in theaters on July 3, 1985. [E1]" cited to a trilogy/SNES/pinball paragraph containing "1985" only as a diegetic time-period reference (`...back to the real 1985...`) and no "July" anywhere. Year-alone substring would have passed; requiring **all** date components asserted in the claim catches it.

- *Entity-list lazy-anchor.* "Name the Simpsons family members and their pets." A model emits a list including a training-prior entity that does not appear in any cited chunk. The any-match semantics passes the claim if at least one of the named entities is anchored — demote-don't-reject keeps the auditor sees the listing intact while the cap-at-HYBRID signal flags the claim as not fully grounded.
- *Count lazy-anchor.* "How many wives did Henry VIII have?" answered "Henry VIII had six wives. [E1]" cited to a span that uses the digit form ("...married 6 women...") or the word form ("...had six wives..."). Digit↔word equivalence ("six" ↔ "6", "twelve" ↔ "12") passes either form; only when neither form appears does count warrant fail. Catches the case where the model emits a count from training prior cited to a span with a *different* number.
- *Why-cause lazy-anchor.* "Why did the Roman Empire fall?" answered with a list of causes (economic, political, military) cited to a span that talks about emperor names and battle dates only. At least one cause anchor (proper noun OR ≥5-char common noun outside a small stopword pool) must appear in some cited span — the any-match semantics tolerates the model adding extra causes from prior as long as at least one is grounded.

The anchor-class generalization landed instead of a per-question-type contract framework (`release_date_lookup` + `box_office_lookup` + `birthday_lookup` etc.). The general primitive — *claim asserts a specific lexical anchor → cited span must contain it* — catches both failure classes deterministically without a typed-rule library to maintain. Typed contracts earn their slot only if bench evidence shows general anchors miss specific cases the lexical primitive cannot reach.

Limitation accepted: ISO date format (`1985-07-03`) does not satisfy the month-name requirement on its own. Real-world Wikipedia prose uses month names ("released on July 3, 1985"), so the trade-off favors catching the common case. If bench shows ISO false-rejections, layer in a month-number alternation to the month-name regex.

The strict no-manual-quotes rule that earlier drafts of this section described was removed in late April 2026. Hermes-class models routinely paraphrase source prose but copy named-quoted phrases verbatim (e.g. `"Constitution State"` from a Connecticut span); rejecting any claim that contained any `"` character was rejecting factually-correct, source-grounded claims for cosmetic punctuation. The citation-coverage threshold (Rule 5) and pointer cap (Rule 6) carry the weight of catching the synthetic-elision and mega-claim failures the old rule was meant to catch, without false-rejecting on legitimate quoted phrases. The architectural property — **the runtime owns quote text, not the model** — is preserved because the renderer still interpolates literal spans from the evidence map at display time; the model never types the quote string.

13.9.3 Spotlight render — cited evidence, one window per claim

The renderer interpolates literal source spans from the evidence map at display time. A naive truncation (`span[:200]`) shows the same article-intro sentence under every claim when the model lazy-anchors at the topic id. Aborist instead extracts content tokens from each claim text (≥4-char alphanumeric tokens, stopword-filtered) & runs a **density rank** over their match positions: find ALL occurrences of ALL tokens in the cited span, then for each candidate position count how many DISTINCT tokens fall within ±half-window. Pick the position with maximum distinct count; tie-break on smallest index for determinism. The displayed window centers on that position, expanded outward to the nearest sentence boundaries (so the excerpt never cuts mid-word). If no token matches, the renderer falls back to the leading sentence(s) with a trailing ellipsis.

Density beats first-match-of-longest-token on a load-bearing failure mode. Pre-density, "who is Homer Simpson's boss?" — answered "Homer Simpson's boss is Mr. Burns. [E1]" — would spotlight the slice where `homer` first matched in the cited chunk, often a voice-actor bio paragraph ("For voicing Homer, Castellaneta has won..."). The full chunk genuinely contained the load-bearing slice ("...Homer is often ignored and completely forgotten by his boss Mr. Burns..."), but the displayed excerpt landed on a non-load-bearing slice and the audit trail looked dishonest. Density rank picks the slice where `homer` + `boss` + `burns` cluster together — the slice that actually justifies the claim.

The pointer is also formatted to disambiguate from the source-rank list rendered alongside. The bracket grew from `[E5: "<excerpt>"]` to:

```
[E5 | <source title> | <chunk_root prefix>: "<excerpt>"]
```

Closes a 2026-05-01 visual confusion observed on an Orwell run where `[E5: "..."]` displayed alongside a "sources (8)" list whose `[5]` slot was a different document — readers parsed `E5` as a 1-indexed source rank when it is actually a chunk pointer. Title comes from the evidence object (URI tail fallback); chunk_root prefix is the first 8 hex chars — enough for the operator to disambiguate while staying compact. *Same evidence pointer, different displayed sentence per claim* — anchored on what the claim is actually about, with the source label visible in the bracket itself.

13.9.4 Per-chunk granularity & query-relevance ordering

A whole-document evidence object lets a small model lazy-anchor: cite `[E1]` (the topic article) for every claim & let the renderer figure out the rest. The fix is structural — each retrieved source contributes one `EvidenceObject` per chunk, not one per source. The model sees 30+ short paragraph-scale spans tagged `E1`…`E30` & must pick the specific paragraph that supports each claim, or omit the claim.

Within each source, chunks are ranked by query-token overlap before `pointer_id` assignment:

```
sort key:
    -distinct_query_tokens_present     # primary: more topic words wins
    -total_mentions                    # tiebreaker: density wins
    +chunk_idx                         # final tiebreaker: doc order
```

So `E1` of a source is the chunk most likely to textually support the question, not the chunk that happens to be first in document order. This is a soft signal — token overlap, never enters the proof path — but it shapes the model's anchoring without giving it semantic discretion. Hermes-3-8B's lazy-anchor habit (cite the first thing you see) lands on a relevant chunk by accident.

13.9.5 9-stage CTI run-DAG

The per-run Merkle-DAG (`run_dag_root`) grows from 7 stages to 9 in pointer mode:

```
quote mode (default, 7):
    question / retrieval / context / prompt / answer /
    verify / final_label

pointer mode (9):
    question / retrieval / evidence_map / prompt / raw_answer /
    parsed_claim_lattice / verify / render / final_label
```

The new stages commit each provenance step independently:

- `evidence_map` — Merkle root over sorted `evidence_ids`, order-independent so two retrieval runs that returned the same chunks in different orders produce the same root.
- `raw_answer` — sha256 of the model's literal pointer-line output.
- `parsed_claim_lattice` — sha256 of the canonical-JSON of the parsed claim nodes (claim_text + content-addressed evidence_ids per claim). Reordering claims changes the hash.
- `render` — sha256 of the rendered prose with literal spans interpolated.

`context` drops out — the context IS the evidence map. `answer` splits into raw / parsed / render so an auditor can reconstruct each transformation step from the persisted blob. Quote mode keeps the original 7-stage shape so pre-G0 `run_dag_root` values stay valid.

13.9.6 No iterative repair — one-shot discipline

Pointer mode skips the mechanical-repair + reprompt-feedback loop that 13.8 reserves for quote mode. The benchmark contract is *zero-shot ideal, one-shot maximum, no repeated retries until success*. A run may fail honestly. A run may self-label uncertainty. A run may downgrade itself. A run may not keep trying until it passes. UNGROUNDED beats hallucinated success.

The hard discipline carries a measurable cost: small models drop the bracket protocol on perhaps half of cold runs. The lever is the prompt, specifically a one-shot worked example showing the format end-to-end. With the example, JP dinosaurs benchmark went 0/13 → 10/11 → 16/17 verified across the iterations that hardened the pipeline. Without the example, runs were near-random in protocol compliance.

13.9.7 What pointer mode does NOT prove

The pointer-mode verifier proves *evidence-linked, role-OK, coverage-met, source-title-relevant,* and (when the active anchor class fires) *anchor-named-in-cited-chunk* for every active class — proper-noun on relation-shape questions, date components on date-asserting claims, named entity on entity-list-shape questions, count token on count-shape questions, cause anchor on why-shape questions. It does **not** prove the cited evidence semantically entails the claim. A model citing `[E1]` for "Brachiosaurus exhibits social herding behavior in the film" passes the eight hard checks if E1's source is the Jurassic Park film article (Rule 8 title-overlap), Brachiosaurus appears in E1 (Rule 5 coverage), and the relation/date/list/count/why anchor classes don't fire on a "behavior" claim — even if E1's text says nothing about social herding. A model citing `[E1]` for "Back to the Future was released on July 3, 1985" passes the date-anchor check if E1 contains both "July" and "1985" — even if those tokens sit next to each other in a different release context (game, ride, sequel) than the original film's theatrical release. The structural-and-lexical floor is what is proved; the semantic ceiling above it is what is not.

That ceiling gap is by design. NLI-grade entailment requires either a textual-entailment model (soft signal) or LLM-as-judge (Hermes round-trip) — both of which would re-introduce the soft/hard boundary leak the architecture is built to prevent. The honest report is "this claim is evidence-linked, role-OK, lexically grounded, and warrant-named; semantic entailment is not yet hard-verified." The spotlight renderer makes residual gaps visible: when the cited span has no token cluster (the density rank finds no high-density region), the displayed excerpt falls back to the leading window — an auditor reading EVIDENCE_LINKED + leading-window excerpts under a relation-shape question immediately sees "the verifier accepted on coverage but the chunk's load-bearing slice for this claim is sparse," even when warrant-lite passed because the answer entity exists somewhere in the chunk.

Warrant-lite (Rule 7) and title-relevance (Rule 8) are the first semantic-grounded checks to enter the hard layer, and they earn the entry by staying lexical. The lattice now covers six layers cleanly: provenance commitment (Merkle-AGI), source linkage (Reverse-RAG), claim decomposition (CTI), anchor-class warrant (Rule 7 — proper-noun, date, entity-list, count, cause anchors), source-title relevance (Rule 8 — claim must share content tokens with the cited source's title, catching retrieval-driven hallucinations where the source itself is structurally unrelated), & quote-safe rendering (density-spotlight). The remaining layer — semantic / NLI-grade entailment — is intentionally deferred until a deterministic-enough entailment substrate exists to enter the proof path without staining the soft/hard boundary. Today's stack is honest about that ceiling: when a small model emits a claim its cited evidence cannot semantically support, the verdict is HYBRID with a smell-sidecar warning, never STRICT-with-bogus-citation.

13.9.8 Claim-count ceiling — TOO_MANY_CLAIMS as defense in depth

The 2026-04-30 york-england run revealed an answer-shape failure orthogonal to per-claim verification: `"tell me all there is to know about york england?"` prompted the model to emit 26-59 claim-pointer pairs of which only 2-4 verified. The model treated the prompt-shape as a license to spam encyclopedic claims from training & assign each a coincidental pointer match. An atomic-claim prompt rule (one focused fact per line, multiple distinct facts become multiple lines) reduced the typical case to ~10 claims, but the runaway pattern persists as a model-behavior risk.

`policy["claim_lattice_max_claims_per_answer"]` caps the number of claims a single answer can emit (default 12). Beyond the cap, `verify_claim_lattice` & `verify_claim_lattice_json` record a `TOO_MANY_CLAIMS` violation that demotes STRICT → HYBRID via the existing violation path. The cap doesn't truncate — every claim still verifies so the operator sees the full evidence of the runaway. 12 admits

typical entity-list questions (5-7 dinosaurs, Simpsons family + pets) while flagging the descriptive-prompt runaway. Folds into `governance_policy_hash` on change.

13.9.9 Label discipline — four-rung ladder for claim-lattice modes

The audit_mode token `STRICT` carries different weight across verifier methods. In quote / span / entity / paraphrase mode, STRICT means "every cited evidence unit textually matches a pinned source span" — a strong claim about lexical grounding. In claim-lattice modes (pointer & JSON), STRICT means "every pointer resolves to a valid evidence object whose source_role is allowed AND every cited span passes citation-coverage." That property is structurally weaker: synthesis-heavy claims (a model joining multiple cited spans into a single explanatory sentence) can pass all hard checks without semantic entailment of the joined assertion.

The renderer maps lattice-mode audit_mode tokens to a four-rung ladder where each rung names a strictly stronger property:

```
POINTER-LINKED       pointer/source/chunk verified;
                     warrant either did not apply or failed
ANCHOR-WARRANTED     pointer-linked + warrant passed where it
                     ran; other soft demotes may apply
EVIDENCE-WARRANTED   anchor-warranted + no soft demotes
UNGROUNDED           no verified pairs
```

HYBRID adds a `-PARTIAL` suffix to whichever rung applies. `ENTAILMENT-VERIFIED` reserves a fifth rung for a future committed entailment engine; today's stack does not produce it.

Rung discrimination uses the existing violations list — no new verifier output field needed. `WARRANT_MISSING` violations drop the rung to POINTER-LINKED (warrant ran & failed for at least one claim). Soft-demote violations like `LAZY_ANCHOR_DEMOTED` / `POINTER_OVERFLOW_TRIMMED` / `TOO_MANY_CLAIMS` / `BARE_NAME_CLAIM` cap the rung at ANCHOR-WARRANTED. No demotes & warrant clean → EVIDENCE-WARRANTED.

Quote / span / entity / paraphrase modes keep their original audit_mode tokens (those verify against pinned spans, not synthesis). The schema column stays `{STRICT, HYBRID, UNGROUNDED}` so the v9.8 invariant holds across cache, audit chain, & mesh wire format. Pure display-layer transformation; cache_key, governance_policy_hash, & all programmatic callers see the original audit_mode unchanged. The earlier two-rung surface (`EVIDENCE-LINKED` / `EVIDENCE-LINKED-PARTIAL`) was the v1 of this work; the four-rung ladder lands the property names directly so operators read what the lexical verifier could confirm without inferring from the verifier_method tail.

13.10 Operator surface — capacity metrics & retrieval keywords

Two operator-facing additions surfaced from bench-driven iteration & live-query feedback. Both stay outside the proof path (never enter `cache_key`) but get displayed in the human render so an operator can reason about the system's behavior without parsing run-DAG blobs.

13.10.1 Capacity metrics — char-level prompt accounting

Per-call result dict gains a `prompt_chars` breakdown:

```
{
  "system_prompt":      <int>,   # system message
  "grounding_reminder": <int>,   # restated rule (claim-lattice modes)
  "user_question":      <int>,   # original question text
  "evidence_or_context":<int>,   # rendered evidence map / wikitext context
  "messages_total":     <int>,   # sum of all messages[].content
}
```

Plus `answer_chars` for the rendered answer length. Char-level: model-agnostic, fast, & a workable proxy for prompt tokens (~4 chars/token English prose, ~2.5 for JSON-evidence-block heavy contexts). The bench harness aggregates these into a *strict-rate by prompt-size bucket* table

(<8KB / 8-16KB / 16-32KB / 32-64KB / >=64KB) so an operator can answer "*does context size correlate with verdict quality?*" at scale.

The CLI human render emits a one-liner under the source list:

```
capacity: prompt 21,996 chars (sys 1,292 + reminder 342 +
          evidence 20,304 + question 30) → answer 247 chars
```

13.10.2 Retrieval keywords — operator-supplied search hints

Long discursive questions sometimes under-retrieve because their content tokens get diluted by template phrasing. A 2026-05-01 case:

```
Q: "what technology may enable one person to reconstruct
    another person's thoughts or ideas without speaking..."
```

returned HYBRID 6/10 with mostly-tangential sources (videoconferencing, telepathy fiction). Appending `"transcranial knowledge acquisition"` to the question lifted the verdict to STRICT 1/1 with the *Neurotechnology* article cited correctly. The user-supplied keywords narrowed OR-mode FTS5 to topical articles.

The `--retrieval-keywords` flag (CLI: `K="..."` on `make query`) makes this pattern explicit. Plumbing: keywords concatenate with the question for the FTS5 search & title-filter token set only. The LLM still sees the original question; the verifier still checks against the original question; `question_hash` still computes from the original text. The 8-dim cache_key changes only via downstream effects on `context_root` & `conversation_hash` (different sources fed to LLM means different conversation state).

This indirect cache_key effect named a provenance gap: keywords changed which sources got chosen, but the keyword string itself entered no Merkle commitment. An audit could recover "these documents were selected" but not "these were the keywords that pulled them in." Two runs with the same question & different keywords that surfaced identical sources were Merkle-indistinguishable.

Ticket #000001 closed that gap 2026-05-02 via `aborist/qa/retrieval_plan.py`:

```
@dataclass(frozen=True)
class RetrievalPlan:
    retrieval_keywords: str
    top_k: int
    over_fetch: int
    max_context_chars: int
    shard_ids: tuple[str, ...]

def retrieval_plan_hash(plan: RetrievalPlan) -> str:
    # SHA-256 over canonical-JSON. Deterministic per call;
    # folds into the run-DAG retrieval stage as a bound input.
```

`aborist/qa/dag.py:build_run_dag` now accepts a `retrieval_plan_hash` parameter. When provided, the retrieval-stage hash binds BOTH the plan (input) & the sources_summary (output) — `H({"retrieval_plan_hash": <hex>, "sources_summary_hash": <hex>})`. When omitted, falls back to the historical sources-summary-only hash so pre-#000001 `run_dag_root` values stay byte-stable. `aborist/qa/query.py` constructs the plan per call & threads the hash through.

Two runs with identical sources but different retrieval keywords now produce different `run_dag_root` values. Audit can recover the keywords from the persisted run-DAG blob. The 8-dim cache_key invariant holds — the plan does NOT enter `question_hash` or `governance_policy_hash`, since keywords stay operator metadata, not user-question intent. SQL-column queryability & `retrieval_plan_built` audit events stay deferred (ergonomic enhancements that earn their own tickets when the run-DAG-blob path proves too friction-heavy).

13.11 Tickets — design log convention

Aborist's design log lives in `docs/` as flat `ticket-NNNNNN-<slug>.md` files. Tickets capture actionable proposals awaiting go/no-go, scoped defects, or queued enhancements. Architecture references (`cti-architecture.md`, `mesh.md`) & bench journals (`qa-modes-bench-2026-04-30.md`) stay un-numbered: tickets propose change & await decision; reference docs describe state of the world.

`docs/TICKETS.md` carries the index plus a `Next ID` line that gets bumped atomically with new ticket commits. Every ticket starts with:

```
# Ticket #NNNNNN — <short title>

**Status:**       open · awaiting go/no-go    (or "in progress", "closed")
**Opened:**       YYYY-MM-DD
**Scope:**        <one-liner>
**Audience:**     <who reads this>
**Hard constraint:** <invariants this ticket commits to NOT break>
```

Closed tickets stay in place as the design log; they don't get deleted.

Tickets shipped as of 2026-05-02:

- **#000001** — Retrieval-keywords audit gap. Closed via `RetrievalPlan` dataclass + `retrieval_plan_hash` folded into the run-DAG retrieval stage (see 13.10.2 above). Two runs with identical sources but different retrieval keywords now produce different `run_dag_root` values.
- **#000002** — Reference-Frame Polarity Contract (Module L). Closed via the frame detector + polarity preamble below (see 13.12).
- **#000003** — Anchor-class warrant generalization. Closed via per-shape detectors + extractors for entity-list / count / why-cause shapes (see 13.9.2 above).
- **#000004** — Directive coverage in bench summary. Closed at landing time (the directive_compliance per-row column in `bench/qa_sweep.py` plus the markdown summary's `directive coverage` table; see 13.13).
- **#000005** — Label ladder migration. Closed via the four-rung POINTER-LINKED → ANCHOR-WARRANTED → EVIDENCE-WARRANTED → UNGROUNDED ladder (see 13.9.9 above), with ENTAILMENT-VERIFIED reserved for a future committed entailment engine.

13.12 Reference-frame polarity contract (Module L)

The phrase-pattern retrieval route (13.4.8) closed the **retrieval** side of reference-frame failure — for an Orwell-shape query like `"has oceania always been at war with east asia"`, the verbatim 5-token sequence `"oceania always been at war"` surfaces the `Nineteen_Eighty-Four` article into top-K. The remaining gap was answer **shape**: the model's pre-#000002 response read `"The text does not directly state..."` — a defensive hedge that ignored the literary frame the retrieved evidence opened.

The substrate now classifies the query's intended frame & nudges the prompt accordingly. New module `aborist/qa/frame.py` adds a `FrameDetection` dataclass plus `detect_frame(question, sources, phrase_match_roots)` heuristic. The classification fires when:

1. The phrase-pattern retrieval route surfaced at least one source (`phrase_match_roots` non-empty).
2. At least one phrase-matched source is a reference work, determined by either:
    - title parenthetical disambig (`"(novel)"`, `"(film)"`, `"(play)"`, `"(franchise)"`, etc.); or
    - body sample contains ≥3 distinct fiction markers from a set including `novel`, `published`, `protagonist`, `plot`, `narrator`, `adaptation`, `directed`, `starring`, `screenplay`, etc.

Distinct-marker count keeps the heuristic robust against single-marker repetition: a history article saying `"a novel approach"` twice doesn't trip the detector; a reference-work lead clusters several distinct markers (`novel` + `published` + `characters` + `plot` + `protagonist`).

The body sample uses the **article LEAD** (chunk_idx=0, post-wikitext-strip) — fiction markers cluster in the lead on Wikipedia, not in plot chunks the query-relevance ranker may have surfaced higher. Wikitext strip ensures markers buried under `[[wikilinks]]` & `{{templates}}` reach the density check.

When `frame_kind == "reference"`, `aborist/qa/query.py` injects a polarity preamble as a user-role message before the existing grounding_reminder:

```
"This question may be a reference to {reference_title}. When you
answer, distinguish what the cited work depicts as actual continuity
from what in-universe propaganda or characters claim within it. Cite
evidence for each substantive claim using the pointer format above."
```

The renderer adds a `reference frame: <title>` line to the human output. Result dict gains a `frame_detection` field carrying the kind / title / URI / confidence — sidecar; never enters cache_key, governance_policy_hash, or the run-DAG.

Live verification on the Orwell case:

```
Pre-#000002:  "The text does not directly state that Oceania has
               always been at war with East Asia."

Post-#000002: "In George Orwell's dystopian novel Nineteen
               Eighty-Four, the nation of Oceania is always at war
               with Eastasia, but this is a result of propaganda &
               doublethink, not actual historical continuity. The
               war with Eastasia is a fabricated conflict to
               maintain control..."
               reference frame: Nineteen Eighty-Four
```

The post-fix answer distinguishes the propaganda claim from the fictional-actual continuity, the multi-frame compilation the polarity contract abstraction proposed. Literal queries (`"what is the capital of france?"`) keep their clean single-frame answers — the polarity preamble injects only when `detect_frame` classifies the query as reference.

Per-frame answer compilation as a runtime concern (programmatically emit per-frame paragraphs rather than nudging via prompt) stays deferred. The prompt-side approach lands first; the runtime compiler earns its slot if bench evidence shows the prompt nudge proves unreliable.

13.13 Engineering doctrine — seven-point program & five-step algorithm

The session that landed #000001-#000005 also formalized the architectural directives that drove them. Two artifacts now sit upstream of every ticket: the **seven-point program** (the directive catalogue) & the **five-step algorithm** (the discipline for evaluating change). Both live in `docs/` & in `CLAUDE.md` so every shift, human or agent, walks past them before touching code.

13.13.1 Seven-point program

Distilled 2026-05-01:

```
Stop making Hermes prove things. Make Hermes emit pointer
clauses; build CTI internally; bind the retrieval map &
evidence map; verify pointers deterministically; add general
anchor-class warrant before semantic NLI; rename labels
honestly; & automate only after these invariants are
test-pinned.
```

Each directive maps to a structural pin in `tests/test_directives.py` & to the run-time substrate that enforces it:

- **D1** — No LLM in hard verifier path. Substrate (`aborist/qa/verify.py` signatures forbid `chat_client` / `llm` / `judge` / `model` params; CHECK constraint enum on `verifier_method` excludes LLM-judge tokens).
- **D2** — Hermes emits pointer clauses. Substrate (`ANSWER_MODES` carries `claim_lattice_pointer` & `claim_lattice`; runtime pointer parser at `aborist/qa/parse_claims.py`).
- **D3** — Build CTI internally. Closed via #000002 (frame detector + polarity preamble + multi-frame answers).
- **D4** — Bind retrieval map AND evidence map. Closed via #000001 (`retrieval_plan_hash` folded into the run-DAG retrieval stage).
- **D5** — Verify pointers deterministically. Substrate (seven hard checks in `verify_claim_lattice`; deterministic byte-for-byte across runs).
- **D6** — Anchor-class warrant before semantic NLI. Closed via #000003 (entity-list / count / why-cause shapes added to the warrant dispatch).
- **D7** — Rename labels honestly. Closed via #000005 (four-rung ladder POINTER-LINKED → ANCHOR-WARRANTED → EVIDENCE-WARRANTED → UNGROUNDED, with ENTAILMENT-VERIFIED reserved).
- **D8** — Automate only after the invariants are test-pinned. Closed via #000004 (bench-side `directive_compliance` per-row column + `directive coverage` markdown summary section).

All seven structural directives close as of 2026-05-02. `docs/seven-point-program.md` carries the catalogue; `tests/test_directives.py` carries 23 anti-regression tests that fail by name when a future PR silently weakens any directive (e.g. adds a `chat_client` parameter to `verify_quotes`, or admits `llm` / `judge` / `nli` to the `verifier_method` CHECK constraint enum).

13.13.2 Five-step algorithm

Walk these in sequence; skipping a step makes the next ones expensive:

```
1. Make the requirements less dumb. Every requirement gets a
   person's name, not a department. If you can't name who asked
   or which defect closed, the requirement is suspect.

2. Delete the part or the process. If you aren't putting back at
   least 10% of what you delete, you aren't deleting hard
   enough. The verifier-stays-binary discipline & "no soft
   signals in the hard chain" rules are deletion-first
   guardrails.

3. Simplify & optimize. Only after 1+2. Don't optimize a
   process that shouldn't exist.

4. Accelerate cycle time. Only after simplifying. Don't dig
   the grave faster.

5. Automate. Last, not first. Hand-rolled before scripted,
   scripted before declarative, declarative before generated.
```

The five-step algorithm shows up structurally in the ticket bodies. Per-class warrant policy gates were proposed in #000003 §3.4 & deferred under step 2: the single `warrant_check_enabled: bool` flag is the minimum-viable mechanism. Per-class flags earn their slot when bench evidence shows over-firing on a specific class — not before. The `verifier_steps_ran` field proposed in #000005 §3 was similarly deferred: the existing `violations` list carries enough signal to discriminate the four ladder rungs without adding a new field.

13.13.3 Bench-as-build-gate

`bench/qa_sweep.py` enforces directive coverage on every sweep. Each row carries a `directive_compliance: dict[str, bool]` field for the directives whose pinning is observable per-row (D2 / D3 / D4 / D6 / D7). The summary markdown emits a `## directive coverage (seven-point program)` table showing per-mode pass rates as `count/N (pct%)`. D1 / D5 / D8 stay system-global properties tracked once in the program doc.

Bench scores that climb without directive coverage climbing are graveyard-digging (five-step algorithm step 4). The bench harness IS the automation substrate, so it must enforce the directives before any feature gets layered on top.

14. Related Work

Graphify (*pip install graphify*) builds navigable knowledge graphs from local folders: code in 13 languages, PDFs, images, & markdown. One command produces an Obsidian vault, a wiki, & a graph queryable in plain English. Graphify's authors report significant token reductions for graph-traversal queries versus raw file reads. Our work complements Graphify: where Graphify builds graphs over static local folders, Merkle Providence describes verifiable answer chains over live web documents & git-versioned codebases, with cryptographic provenance for every cached result. No prototype of our integration exists yet.

Git / Mercurial object model. Both git & Mercurial implement content-addressed Merkle DAGs natively. Git's commit object hashes our entire tree; Merkle Providence inherits this as a free cache boundary signal. Our work makes this implicit property explicit & useful for ML inference caching.

Traditional RAG (Lewis et al., 2020). Retrieval Augmented Generation introduced server-side vector retrieval for language models. Our Reverse RAG paper documented client-side inversion of this pipeline. Merkle Providence extends Reverse RAG with verifiable caching, eliminating our repetition tax that RAG systems do not address.

Zero-Knowledge Proofs. Merkle membership proofs constitute a subset of ZKP: proving a leaf exists in a tree without revealing other leaves. Our work applies this primitive to ML answer provenance without requiring full ZKP apparatus (zk-SNARKs, Bulletproofs). We trade completeness for simplicity & speed.

Acknowledgements

Our Merkle cache idea emerged from a conversation on X Spaces. We thank **@AR_Milne** & **@without_rulers** for a discussion that sparked the core insight: that a Merkle tree over document content could serve as a fast, verifiable cache layer for machine learning inference. Ideas born in public conversation belong to our commons.

Citation

```
Russell Ballestrini, David Wong, Riley Morgan.
"Unturf Automated General Intelligence: Merkle Providence Reverse RAG."
unfirehose.com, 2026.
https://unfirehose.com/merkle-providence-reverse-rag.html
```

License

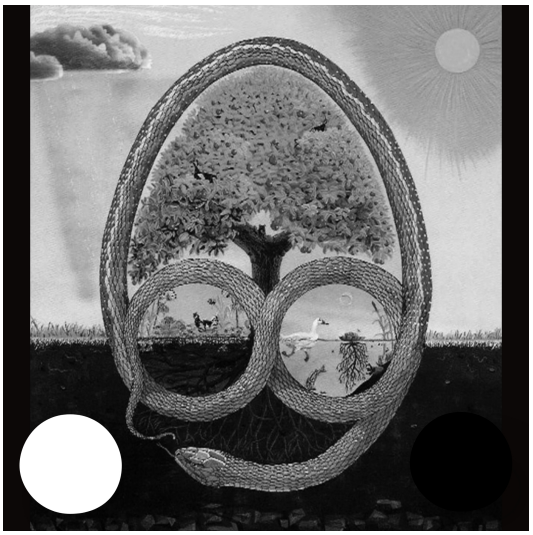

GNU Affero General Public License v3

```
AGPL-3.0-only

PERMACOMPUTER PREAMBLE - NO WARRANTY

This is free software for the public good of a permacomputer hosted at
permacomputer.com - an always-on computer by the people, for the people. One
which is durable, easy to repair, and distributed like tap water for machine
learning intelligence.

The permacomputer is community-owned infrastructure optimized around four values:

  TRUTH    - Source code must be open source & freely distributed
  FREEDOM  - Voluntary participation without corporate control
  HARMONY  - Systems operating with minimal waste that self-renew
  LOVE     - Individual rights protected while fostering cooperation

This paper contributes to that vision by documenting Merkle Providence Reverse
RAG: a provenance-preserving cache layer that lets small language models punch
far outside their parameter class. Code is seeds to sprout on any abandoned
technology.

Learn more: https://www.permacomputer.com

NO WARRANTY. THE SOFTWARE IS PROVIDED "AS IS" WITHOUT WARRANTY OF ANY KIND.
```

Copyright (C) 2026 Russell Ballestrini & David Wong & Riley Morgan
russell@unturf.com
david.s.wong@pm.me
guybriley02@gmail.com
https://www.unturf.com/software
https://www.permacomputer.com
https://uncloseai.com
https://unfirehose.com
https://russell.ballestrini.net

GNU AFFERO GENERAL PUBLIC LICENSE
Version 3, 19 November 2007

Copyright (C) 2007 Free Software Foundation, Inc. <https://fsf.org/>
Everyone is permitted to copy and distribute verbatim copies
of this license document, but changing it is not allowed.

Preamble

The GNU Affero General Public License is a free, copyleft license for software and other kinds of works, specifically designed to ensure cooperation with the community in the case of network server software.

The licenses for most software and other practical works are designed to take away your freedom to share and change the works. By contrast, our General Public Licenses are intended to guarantee your freedom to share and change all versions of a program--to make sure it remains free software for all its users.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for them if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs, and that you know you can do these things.

Developers that use our General Public Licenses protect your rights with two steps: (1) assert copyright on the software, and (2) offer you this License which gives you legal permission to copy, distribute and/or modify the software.

A secondary benefit of defending all users' freedom is that improvements made in alternate versions of the program, if they receive widespread use, become available for other developers to incorporate. Many developers of free software are heartened and encouraged by the resulting cooperation. However, in the case of software used on network servers, this result may fail to come about. The GNU General Public License permits making a modified version and letting the public access it on a server without ever releasing its source code to the public.

The GNU Affero General Public License is designed specifically to ensure that, in such cases, the modified source code becomes available to the community. It requires the operator of a network server to provide the source code of the modified version running there to the users of that server. Therefore, public use of a modified version, on a publicly accessible server, gives the public access to the source code of the modified version.

An older license, called the Affero General Public License and published by Affero, was designed to accomplish similar goals. This is a different license, not a version of the Affero GPL, but Affero has released a new version of the Affero GPL which permits relicensing under this license.

The precise terms and conditions for copying, distribution and modification follow.

TERMS AND CONDITIONS

0. Definitions.

"This License" refers to version 3 of the GNU Affero General Public License.

"Copyright" also means copyright-like laws that apply to other kinds of works, such as semiconductor masks.

"The Program" refers to any copyrightable work licensed under this License. Each licensee is addressed as "you". "Licensees" and "recipients" may be individuals or organizations.

To "modify" a work means to copy from or adapt all or part of the work in a fashion requiring copyright permission, other than the making of an exact copy. The resulting work is called a "modified version" of the earlier work or a work "based on" the earlier work.

A "covered work" means either the unmodified Program or a work based on the Program.

To "propagate" a work means to do anything with it that, without permission, would make you directly or secondarily liable for infringement under applicable copyright law, except executing it on a computer or modifying a private copy. Propagation includes copying, distribution (with or without modification), making available to the public, and in some countries other activities as well.

To "convey" a work means any kind of propagation that enables other parties to make or receive copies. Mere interaction with a user through a computer network, with no transfer of a copy, is not conveying.

An interactive user interface displays "Appropriate Legal Notices" to the extent that it includes a convenient and prominently visible feature that (1) displays an appropriate copyright notice, and (2) tells the user that there is no warranty for the work (except to the extent that warranties are provided), that licensees may convey the work under this License, and how to view a copy of this License. If the interface presents a list of user commands or options, such as a menu, a prominent item in the list meets this criterion.

1. Source Code.

The "source code" for a work means the preferred form of the work for making modifications to it. "Object code" means any non-source form of a work.

A "Standard Interface" means an interface that either is an official standard defined by a recognized standards body, or, in the case of interfaces specified for a particular programming language, one that is widely used among developers working in that language.

The "System Libraries" of an executable work include anything, other than the work as a whole, that (a) is included in the normal form of packaging a Major Component, but which is not part of that Major Component, and (b) serves only to enable use of the work with that Major Component, or to implement a Standard Interface for which an implementation is available to the public in source code form. A "Major Component", in this context, means a major essential component (kernel, window system, and so on) of the specific operating system (if any) on which the executable work runs, or a compiler used to produce the work, or an object code interpreter used to run it.

The "Corresponding Source" for a work in object code form means all the source code needed to generate, install, and (for an executable work) run the object code and to modify the work, including scripts to control those activities. However, it does not include the work's System Libraries, or general-purpose tools or generally available free programs which are used unmodified in performing those activities but which are not part of the work. For example, Corresponding Source includes interface definition files associated with source files for

the work, and the source code for shared libraries and dynamically linked subprograms that the work is specifically designed to require, such as by intimate data communication or control flow between those subprograms and other parts of the work.

The Corresponding Source need not include anything that users can regenerate automatically from other parts of the Corresponding Source.

The Corresponding Source for a work in source code form is that same work.

2. Basic Permissions.

All rights granted under this License are granted for the term of copyright on the Program, and are irrevocable provided the stated conditions are met. This License explicitly affirms your unlimited permission to run the unmodified Program. The output from running a covered work is covered by this License only if the output, given its content, constitutes a covered work. This License acknowledges your rights of fair use or other equivalent, as provided by copyright law.

You may make, run and propagate covered works that you do not convey, without conditions so long as your license otherwise remains in force. You may convey covered works to others for the sole purpose of having them make modifications exclusively for you, or provide you with facilities for running those works, provided that you comply with the terms of this License in conveying all material for which you do not control copyright. Those thus making or running the covered works for you must do so exclusively on your behalf, under your direction and control, on terms that prohibit them from making any copies of your copyrighted material outside their relationship with you.

Conveying under any other circumstances is permitted solely under the conditions stated below. Sublicensing is not allowed; section 10 makes it unnecessary.

3. Protecting Users' Legal Rights From Anti-Circumvention Law.

No covered work shall be deemed part of an effective technological measure under any applicable law fulfilling obligations under article 11 of the WIPO copyright treaty adopted on 20 December 1996, or similar laws prohibiting or restricting circumvention of such measures.

When you convey a covered work, you waive any legal power to forbid circumvention of technological measures to the extent such circumvention is effected by exercising rights under this License with respect to the covered work, and you disclaim any intention to limit operation or modification of the work as a means of enforcing, against the work's users, your or third parties' legal rights to forbid circumvention of technological measures.

4. Conveying Verbatim Copies.

You may convey verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice; keep intact all notices stating that this License and any non-permissive terms added in accord with section 7 apply to the code; keep intact all notices of the absence of any warranty; and give all recipients a copy of this License along with the Program.

You may charge any price or no price for each copy that you convey, and you may offer support or warranty protection for a fee.

5. Conveying Modified Source Versions.

You may convey a work based on the Program, or the modifications to produce it from the Program, in the form of source code under the terms of section 4, provided that you also meet all of these conditions:

a) The work must carry prominent notices stating that you modified it, and giving a relevant date.

b) The work must carry prominent notices stating that it is released under this License and any conditions added under section 7. This requirement modifies the requirement in section 4 to "keep intact all notices".

c) You must license the entire work, as a whole, under this License to anyone who comes into possession of a copy. This License will therefore apply, along with any applicable section 7 additional terms, to the whole of the work, and all its parts, regardless of how they are packaged. This License gives no permission to license the work in any other way, but it does not invalidate such permission if you have separately received it.

d) If the work has interactive user interfaces, each must display Appropriate Legal Notices; however, if the Program has interactive interfaces that do not display Appropriate Legal Notices, your work need not make them do so.

A compilation of a covered work with other separate and independent works, which are not by their nature extensions of the covered work, and which are not combined with it such as to form a larger program, in or on a volume of a storage or distribution medium, is called an "aggregate" if the compilation and its resulting copyright are not used to limit the access or legal rights of the compilation's users beyond what the individual works permit. Inclusion of a covered work in an aggregate does not cause this License to apply to the other parts of the aggregate.

6. Conveying Non-Source Forms.

You may convey a covered work in object code form under the terms of sections 4 and 5, provided that you also convey the machine-readable Corresponding Source under the terms of this License, in one of these ways:

a) Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by the Corresponding Source fixed on a durable physical medium customarily used for software interchange.

b) Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by a written offer, valid for at least three years and valid for as long as you offer spare parts or customer support for that product model, to give anyone who possesses the object code either (1) a copy of the Corresponding Source for all the software in the product that is covered by this License, on a durable physical medium customarily used for software interchange, for a price no more than your reasonable cost of physically performing this conveying of source, or (2) access to copy the Corresponding Source from a network server at no charge.

c) Convey individual copies of the object code with a copy of the written offer to provide the Corresponding Source. This alternative is allowed only occasionally and noncommercially, and only if you received the object code with such an offer, in accord with subsection 6b.

d) Convey the object code by offering access from a designated place (gratis or for a charge), and offer equivalent access to the Corresponding Source in the same way through the same place at no further charge. You need not require recipients to copy the Corresponding Source along with the object code. If the place to copy the object code is a network server, the Corresponding Source may be on a different server (operated by you or a third party) that supports equivalent copying facilities, provided you maintain clear directions next to the object code saying where to find the

Corresponding Source. Regardless of what server hosts the Corresponding Source, you remain obligated to ensure that it is available for as long as needed to satisfy these requirements.

e) Convey the object code using peer-to-peer transmission, provided you inform other peers where the object code and Corresponding Source of the work are being offered to the general public at no charge under subsection 6d.

A separable portion of the object code, whose source code is excluded from the Corresponding Source as a System Library, need not be included in conveying the object code work.

A "User Product" is either (1) a "consumer product", which means any tangible personal property which is normally used for personal, family, or household purposes, or (2) anything designed or sold for incorporation into a dwelling. In determining whether a product is a consumer product, doubtful cases shall be resolved in favor of coverage. For a particular product received by a particular user, "normally used" refers to a typical or common use of that class of product, regardless of the status of the particular user or of the way in which the particular user actually uses, or expects or is expected to use, the product. A product is a consumer product regardless of whether the product has substantial commercial, industrial or non-consumer uses, unless such uses represent the only significant mode of use of the product.

"Installation Information" for a User Product means any methods, procedures, authorization keys, or other information required to install and execute modified versions of a covered work in that User Product from a modified version of its Corresponding Source. The information must suffice to ensure that the continued functioning of the modified object code is in no case prevented or interfered with solely because modification has been made.

If you convey an object code work under this section in, or with, or specifically for use in, a User Product, and the conveying occurs as part of a transaction in which the right of possession and use of the User Product is transferred to the recipient in perpetuity or for a fixed term (regardless of how the transaction is characterized), the Corresponding Source conveyed under this section must be accompanied by the Installation Information. But this requirement does not apply if neither you nor any third party retains the ability to install modified object code on the User Product (for example, the work has been installed in ROM).

The requirement to provide Installation Information does not include a requirement to continue to provide support service, warranty, or updates for a work that has been modified or installed by the recipient, or for the User Product in which it has been modified or installed. Access to a network may be denied when the modification itself materially and adversely affects the operation of the network or violates the rules and protocols for communication across the network.

Corresponding Source conveyed, and Installation Information provided, in accord with this section must be in a format that is publicly documented (and with an implementation available to the public in source code form), and must require no special password or key for unpacking, reading or copying.

7. Additional Terms.

"Additional permissions" are terms that supplement the terms of this License by making exceptions from one or more of its conditions. Additional permissions that are applicable to the entire Program shall be treated as though they were included in this License, to the extent that they are valid under applicable law. If additional permissions apply only to part of the Program, that part may be used separately under those permissions, but the entire Program remains governed by this License without regard to the additional permissions.

When you convey a copy of a covered work, you may at your option

remove any additional permissions from that copy, or from any part of it. (Additional permissions may be written to require their own removal in certain cases when you modify the work.) You may place additional permissions on material, added by you to a covered work, for which you have or can give appropriate copyright permission.

Notwithstanding any other provision of this License, for material you add to a covered work, you may (if authorized by the copyright holders of that material) supplement the terms of this License with terms:

a) Disclaiming warranty or limiting liability differently from the terms of sections 15 and 16 of this License; or

b) Requiring preservation of specified reasonable legal notices or author attributions in that material or in the Appropriate Legal Notices displayed by works containing it; or

c) Prohibiting misrepresentation of the origin of that material, or requiring that modified versions of such material be marked in reasonable ways as different from the original version; or

d) Limiting the use for publicity purposes of names of licensors or authors of the material; or

e) Declining to grant rights under trademark law for use of some trade names, trademarks, or service marks; or

f) Requiring indemnification of licensors and authors of that material by anyone who conveys the material (or modified versions of it) with contractual assumptions of liability to the recipient, for any liability that these contractual assumptions directly impose on those licensors and authors.

All other non-permissive additional terms are considered "further restrictions" within the meaning of section 10. If the Program as you received it, or any part of it, contains a notice stating that it is governed by this License along with a term that is a further restriction, you may remove that term. If a license document contains a further restriction but permits relicensing or conveying under this License, you may add to a covered work material governed by the terms of that license document, provided that the further restriction does not survive such relicensing or conveying.

If you add terms to a covered work in accord with this section, you must place, in the relevant source files, a statement of the additional terms that apply to those files, or a notice indicating where to find the applicable terms.

Additional terms, permissive or non-permissive, may be stated in the form of a separately written license, or stated as exceptions; the above requirements apply either way.

8. Termination.

You may not propagate or modify a covered work except as expressly provided under this License. Any attempt otherwise to propagate or modify it is void, and will automatically terminate your rights under this License (including any patent licenses granted under the third paragraph of section 11).

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that

copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or rights from you under this License. If your rights have been terminated and not permanently reinstated, you do not qualify to receive new licenses for the same material under section 10.

9. Acceptance Not Required for Having Copies.

You are not required to accept this License in order to receive or run a copy of the Program. Ancillary propagation of a covered work occurring solely as a consequence of using peer-to-peer transmission to receive a copy likewise does not require acceptance. However, nothing other than this License grants you permission to propagate or modify any covered work. These actions infringe copyright if you do not accept this License. Therefore, by modifying or propagating a covered work, you indicate your acceptance of this License to do so.

10. Automatic Licensing of Downstream Recipients.

Each time you convey a covered work, the recipient automatically receives a license from the original licensors, to run, modify and propagate that work, subject to this License. You are not responsible for enforcing compliance by third parties with this License.

An "entity transaction" is a transaction transferring control of an organization, or substantially all assets of one, or subdividing an organization, or merging organizations. If propagation of a covered work results from an entity transaction, each party to that transaction who receives a copy of the work also receives whatever licenses to the work the party's predecessor in interest had or could give under the previous paragraph, plus a right to possession of the Corresponding Source of the work from the predecessor in interest, if the predecessor has it or can get it with reasonable efforts.

You may not impose any further restrictions on the exercise of the rights granted or affirmed under this License. For example, you may not impose a license fee, royalty, or other charge for exercise of rights granted under this License, and you may not initiate litigation (including a cross-claim or counterclaim in a lawsuit) alleging that any patent claim is infringed by making, using, selling, offering for sale, or importing the Program or any portion of it.

11. Patents.

A "contributor" is a copyright holder who authorizes use under this License of the Program or a work on which the Program is based. The work thus licensed is called the contributor's "contributor version".

A contributor's "essential patent claims" are all patent claims owned or controlled by the contributor, whether already acquired or hereafter acquired, that would be infringed by some manner, permitted by this License, of making, using, or selling its contributor version, but do not include claims that would be infringed only as a consequence of further modification of the contributor version. For purposes of this definition, "control" includes the right to grant patent sublicenses in a manner consistent with the requirements of this License.

Each contributor grants you a non-exclusive, worldwide, royalty-free patent license under the contributor's essential patent claims, to make, use, sell, offer for sale, import and otherwise run, modify and propagate the contents of its contributor version.

In the following three paragraphs, a "patent license" is any express agreement or commitment, however denominated, not to enforce a patent (such as an express permission to practice a patent or covenant not to sue for patent infringement). To "grant" such a patent license to a party means to make such an agreement or commitment not to enforce a

patent against the party.

If you convey a covered work, knowingly relying on a patent license, and the Corresponding Source of the work is not available for anyone to copy, free of charge and under the terms of this License, through a publicly available network server or other readily accessible means, then you must either (1) cause the Corresponding Source to be so available, or (2) arrange to deprive yourself of the benefit of the patent license for this particular work, or (3) arrange, in a manner consistent with the requirements of this License, to extend the patent license to downstream recipients. "Knowingly relying" means you have actual knowledge that, but for the patent license, your conveying the covered work in a country, or your recipient's use of the covered work in a country, would infringe one or more identifiable patents in that country that you have reason to believe are valid.

If, pursuant to or in connection with a single transaction or arrangement, you convey, or propagate by procuring conveyance of, a covered work, and grant a patent license to some of the parties receiving the covered work authorizing them to use, propagate, modify or convey a specific copy of the covered work, then the patent license you grant is automatically extended to all recipients of the covered work and works based on it.

A patent license is "discriminatory" if it does not include within the scope of its coverage, prohibits the exercise of, or is conditioned on the non-exercise of one or more of the rights that are specifically granted under this License. You may not convey a covered work if you are a party to an arrangement with a third party that is in the business of distributing software, under which you make payment to the third party based on the extent of your activity of conveying the work, and under which the third party grants, to any of the parties who would receive the covered work from you, a discriminatory patent license (a) in connection with copies of the covered work conveyed by you (or copies made from those copies), or (b) primarily for and in connection with specific products or compilations that contain the covered work, unless you entered into that arrangement, or that patent license was granted, prior to 28 March 2007.

Nothing in this License shall be construed as excluding or limiting any implied license or other defenses to infringement that may otherwise be available to you under applicable patent law.

12. No Surrender of Others' Freedom.

If conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot convey a covered work so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not convey it at all. For example, if you agree to terms that obligate you to collect a royalty for further conveying from those to whom you convey the Program, the only way you could satisfy both those terms and this License would be to refrain entirely from conveying the Program.

13. Remote Network Interaction; Use with the GNU General Public License.

Notwithstanding any other provision of this License, if you modify the Program, your modified version must prominently offer all users interacting with it remotely through a computer network (if your version supports such interaction) an opportunity to receive the Corresponding Source of your version by providing access to the Corresponding Source from a network server at no charge, through some standard or customary means of facilitating copying of software. This Corresponding Source shall include the Corresponding Source for any work covered by version 3 of the GNU General Public License that is incorporated pursuant to the following paragraph.

Notwithstanding any other provision of this License, you have permission to link or combine any covered work with a work licensed under version 3 of the GNU General Public License into a single

combined work, and to convey the resulting work. The terms of this License will continue to apply to the part which is the covered work, but the work with which it is combined will remain governed by version 3 of the GNU General Public License.

14. Revised Versions of this License.

The Free Software Foundation may publish revised and/or new versions of the GNU Affero General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies that a certain numbered version of the GNU Affero General Public License "or any later version" applies to it, you have the option of following the terms and conditions either of that numbered version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of the GNU Affero General Public License, you may choose any version ever published by the Free Software Foundation.

If the Program specifies that a proxy can decide which future versions of the GNU Affero General Public License can be used, that proxy's public statement of acceptance of a version permanently authorizes you to choose that version for the Program.

Later license versions may give you additional or different permissions. However, no additional obligations are imposed on any author or copyright holder as a result of your choosing to follow a later version.

15. Disclaimer of Warranty.

THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

16. Limitation of Liability.

IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MODIFIES AND/OR CONVEYS THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

17. Interpretation of Sections 15 and 16.

If the disclaimer of warranty and limitation of liability provided above cannot be given local legal effect according to their terms, reviewing courts shall apply local law that most closely approximates an absolute waiver of all civil liability in connection with the Program, unless a warranty or assumption of liability accompanies a copy of the Program in return for a fee.

END OF TERMS AND CONDITIONS

How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

```
  To do so, attach the following notices to the program.  It is safest
to attach them to the start of each source file to most effectively
state the exclusion of warranty; and each file should have at least
the "copyright" line and a pointer to where the full notice is found.

    <one line to give the program's name and a brief idea of what it does.>
    Copyright (C) <year>  <name of author>

    This program is free software: you can redistribute it and/or modify
    it under the terms of the GNU Affero General Public License as published by
    the Free Software Foundation, either version 3 of the License, or
    (at your option) any later version.

    This program is distributed in the hope that it will be useful,
    but WITHOUT ANY WARRANTY; without even the implied warranty of
    MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.  See the
    GNU Affero General Public License for more details.

    You should have received a copy of the GNU Affero General Public License
    along with this program.  If not, see <https://www.gnu.org/licenses/>.

Also add information on how to contact you by electronic and paper mail.

  If your software can interact with users remotely through a computer
network, you should also make sure that it provides a way for users to
get its source.  For example, if your program is a web application, its
interface could display a "Source" link that leads users to an archive
of the code.  There are many ways you could offer source, and different
solutions will be better for different programs; see section 13 for the
specific requirements.

  You should also get your employer (if you work as a programmer) or school,
if any, to sign a "copyright disclaimer" for the program, if necessary.
For more information on this, and how to apply and follow the GNU AGPL, see
<https://www.gnu.org/licenses/>.
```